

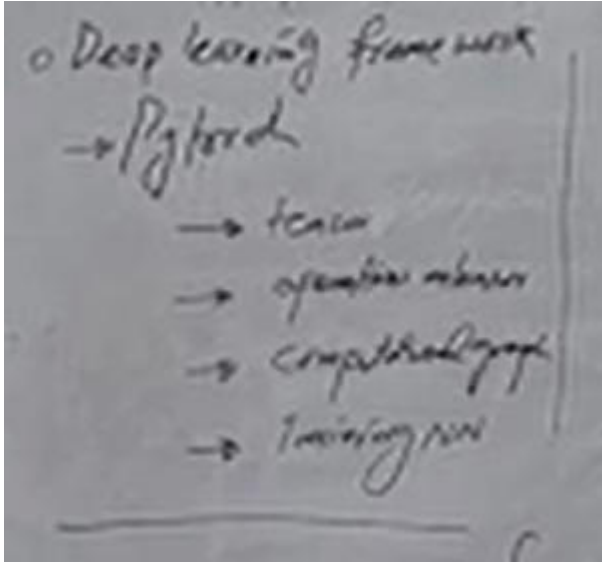
CAP5415

Computer Vision

Yogesh S Rawat

yogesh@ucf.edu

HEC-241



PyTorch Tutorial - I

Lecture 8

Deep learning libraries

- Torch (Lua):
 - <http://torch.ch/>
- PyTorch (Python)
 - <http://pytorch.org/>
- TensorFlow (Python and C++):
 - <https://www.tensorflow.org/>
- Theano (Python)
 - <http://deeplearning.net/software/theano/>
- Keras
 - <https://keras.io/>



PyTorch Tensor



- Similar to NumPy arrays
- They can also be used on a GPU
 - Faster computation
- Random matrix

```
import torch
```

```
x=torch.rand(2,3)
```

```
y=torch.rand(3,3)
```

```
print x
```

```
print y
```

PyTorch Tensor

- Similar to NumPy arrays
- They can also be used on a GPU
 - Faster computation
- All zeros
- Directly from data
- Size of a tensor

```
import torch
```

```
x = torch.zeros(5, 3)
```

```
x = torch.tensor([5.5, 3])
```

```
print x.size()
```

Operations

- Adding tensors
- Indexing

```
x = torch.randn(4, 4)
```

```
y = torch.randn(4, 4)
```

```
print(torch.add(x, y))
```

```
print(x[:, 1])
```

Operations

- Resizing
 - If you want to resize/reshape tensor



```
x = torch.randn(4, 4)
```

```
y = x.view(16)
```

```
z = x.view(-1, 8)
```

```
print(x.size(), y.size(),  
      z.size())
```

Output:

```
torch.Size([4, 4])
```

```
torch.Size([16])
```

```
torch.Size([2, 8])
```

Pop Quiz

Send private message.

Message **to all** will not be considered!

Reshaping tensor

```
x = torch.randn(1, 4, 32, 24)
```

```
y = x.view(8, 2, -1, 3, 8)
```

```
print(y.size())
```

Output shape? 30 seconds!

Torch tensor vs NumPy array

- NumPy array
 - CPU
- Torch tensor
 - GPU

```
a = torch.ones(5)
tensor([1., 1., 1., 1., 1.])
```

```
b = a.numpy()
```

```
a = numpy.ones(5)
b = torch.from_numpy(a)
```



Matrix Multiplication in PyTorch

```
import torch

mat1=torch.randn(2,3)
mat2=torch.randn(3,3)
res=torch.mm(mat1,mat2)

print res.size()
```

Output:

(2, 3)

Batch Matrix Multiplication in PyTorch

```
import torch
```

```
batch1=torch.randn(10,3,4)
```

```
batch2=torch.randn(10,4,5)
```

```
res=torch.bmm(batch1,batch2)
```

```
print res.size()
```

Output:

```
(10L, 3L, 5L)
```



Many Tensor operations in PyTorch...

`torch.mm`

- Matrix multiplication

`torch.bmm`

- Batch matrix multiplication

`torch.cat`

- Tensor Concatenation

`torch.squeeze/torch.unsqueeze`

- Change Tensor dimensions

...

Check documentation at <http://pytorch.org/docs/master/torch.html#tensors>

Computational Graphs


```
import torch
```

```
x = torch.ones(2,2)
```


```
y = torch.ones(2,1)
```

```
w = torch.randn(2,1,requires_grad=True)
```

```
b = torch.randn(1,requires_grad=True)
```



Computational Graphs



```
p = torch.sigmoid(torch.mm(x, w) + b)
# prediction
```

```
loss = -y*torch.log(p) - (1-y) *torch.log(1-p)
# cross-entropy loss
```

```
cost = loss.mean()
# the cost to minimize
```

Automatic Gradient Computation

```
p = torch.sigmoid(torch.mm(x, w) + b)
loss = -y*torch.log(p) - (1-y)*torch.log(1-p)
cost = loss.mean()
```

```
cost.backward()
```

```
print w.grad
print b.grad
```


Training procedure

- Define the neural network
- Iterate over a dataset of inputs
- Process input through the network
- Compute the loss
- Propagate gradients back into the network's parameters
- Update the weights of the network

Training procedure

- Define the neural network
- Iterate over a dataset of inputs
- Process input through the network
- Compute the loss
- Propagate gradients back into the network's parameters
- Update the weights of the network

Build Neural Networks using PyTorch

Neural networks can be constructed using the `torch.nn` package.

Forward

- An `nn.Module` contains layers, and
- A method `forward(input)` that returns the output
- You can use any of the Tensor operations in the forward function

Backward

- `nn` depends on `autograd`
- You just have to define the forward function

Define a Network Class

```
import torch
import torch.nn as nn


class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # create layers

    def forward(self, x):
        # define feed-forward function
```

You don't need to define a backward function!

CNN for MNIST: A Full Example



Example from http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html

Define a CNN Network

```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

Define a CNN Network

```
def __init__(self):  
    super(Net, self).__init__()  
    # 1 input image channel, 6 output channels,  
3x3 square convolution kernel  
    self.conv1 = nn.Conv2d(1, 6, 3)  
    self.conv2 = nn.Conv2d(6, 16, 3)  
    # an affine operation:  $y = Wx + b$   
    self.fc1 = nn.Linear(16 * 6 * 6, 120) # 6*6  
from image dimension  
    self.fc2 = nn.Linear(120, 84)  
    self.fc3 = nn.Linear(84, 10)
```

Define a CNN Network

```
def forward(self, x):  
    # Max pooling over a (2, 2) window  
    x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))  
    # If the size is a square you can only specify a  
single number  
    x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
    x = x.view(-1, self.num_flat_features(x))  
    x = F.relu(self.fc1(x))  
    x = F.relu(self.fc2(x))  
    x = self.fc3(x)  
  
    return x
```


Define a CNN Network

```
def num_flat_features(self, x):  
    size = x.size()[1:] # all dimensions except the  
batch dimension  
    num_features = 1  
    for s in size:  
        num_features *= s  
  
    return num_features
```

Training procedure

- Define the neural network
- Iterate over a dataset of inputs
- Process input through the network
- Compute the loss
- Propagate gradients back into the network's parameters
- Update the weights of the network

Data

- For images
 - Pillow, OpenCV are useful
- For audio
 - Scipy and librosa
- For text
 - NLTK and SpaCy are useful
- Load data into memory as NumPy array
 - Then convert to tensor for GPU

Loading data - torchvision

- Torchvision
 - it's extremely easy to load existing datasets.

```
import torchvision  
import torchvision.transforms as transforms
```



Loading data - torchvision

```
import torchvision
import torchvision.transforms as transforms

transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data',
train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset,
batch_size=4, shuffle=True, num_workers=2)
```

Loading data - torchvision

```
import torchvision
import torchvision.transforms as transforms

transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

testset = torchvision.datasets.CIFAR10(root='./data',
                                         train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset,
                                          batch_size=4, shuffle=False, num_workers=2)
```

Questions?

Sources for this lecture include materials from pytorch.org