

Resources:

- Dr Rawat (UCF)
- Ms. Adeela (PhD scholar PUCIT)
- Dr Shree Nayar (Columbia University)

Lecture 3

Outline

- Image digitization
 - Sampling
 - Quantization
- Image as a function
- Image processing
 - Histogram
 - Histogram equalization

Need of digitization

- Every scene around us is a continuous image that can be represented by infinite number of image points (resolution)
- Each such image points may contain infinitely many possible intensities
- That needing an infinite number of bits for storage
- Theory of real number: Between any 2 given points there infinite number of points
- Obviously such a representation is not possible in any digital computer

Digitization

Issues with continuous image

- Infinite sizes values between two real numbers
- Infinite intensity values between two real numbers
- The process of transforming continuous space into discrete space is called **digitization**
- Vision algorithms use a discrete form of the images

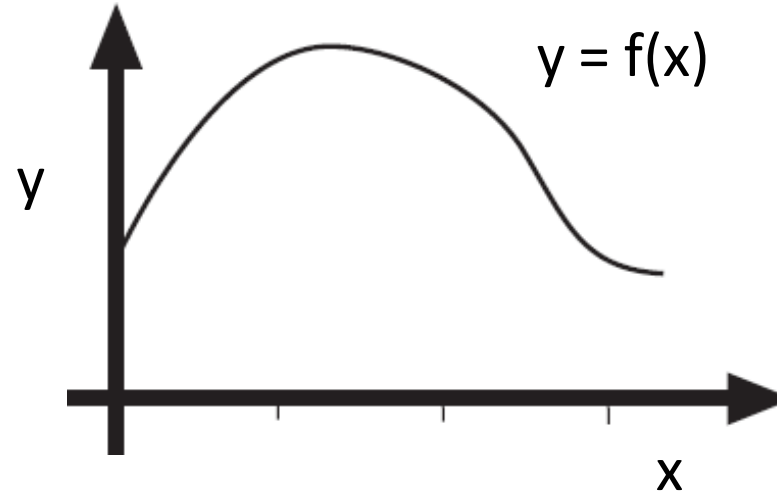


Digitization

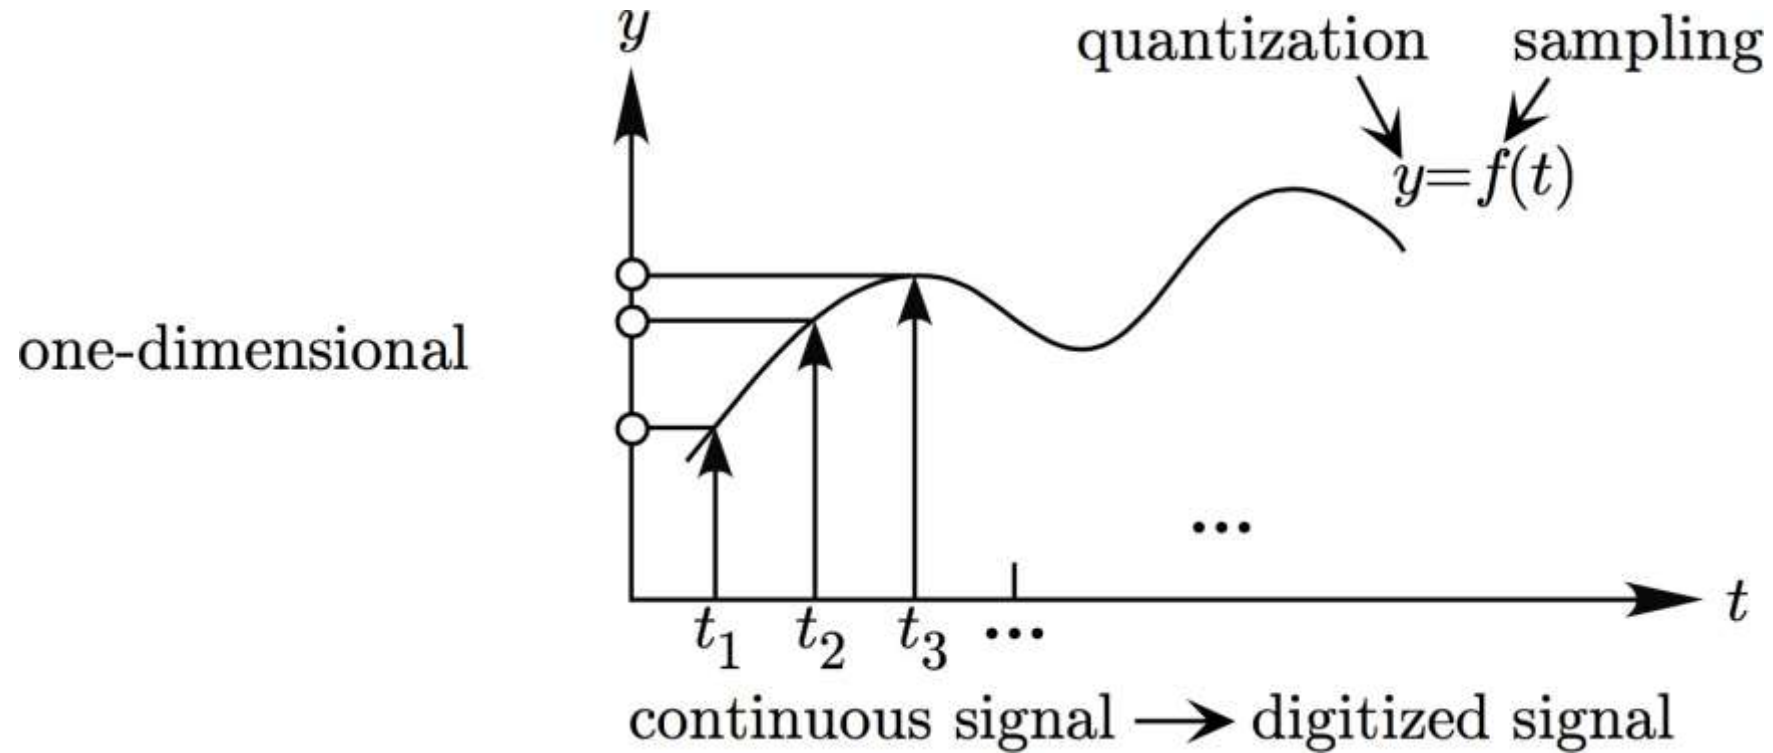
- Function

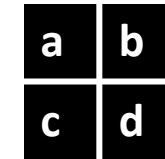
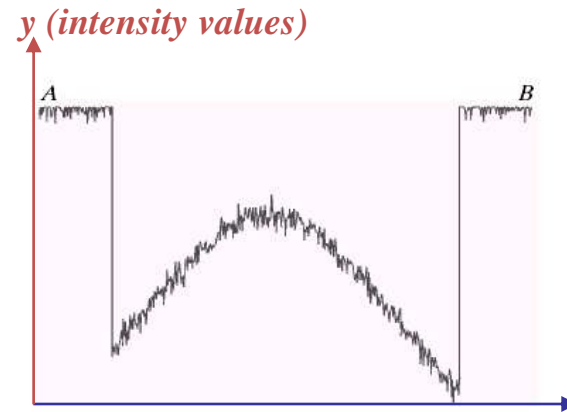
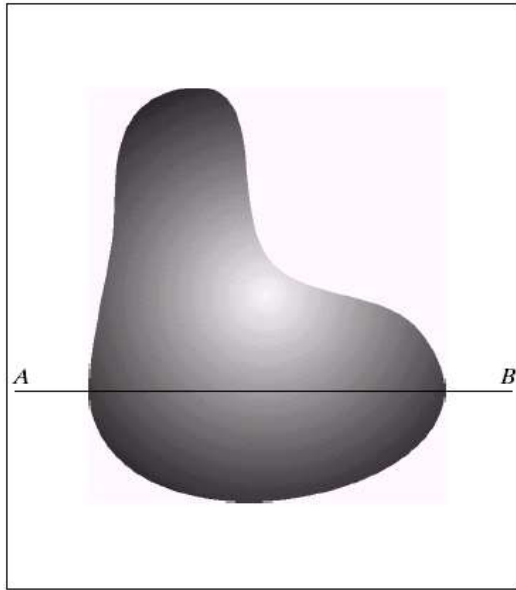
$$y = f(x)$$

- Domain of a function(x)
- Range of a function ($y=f(x)$)
- Sampling
 - Discretization of domain
- Quantization
 - Discretization of range



Digitization of 1D function





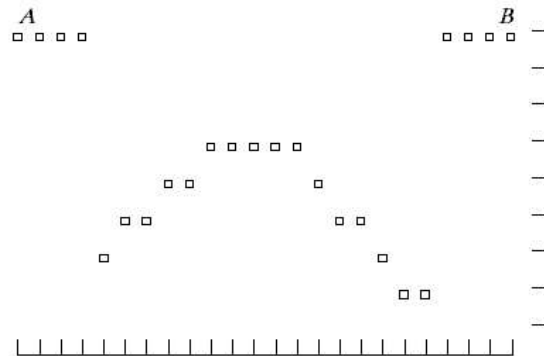
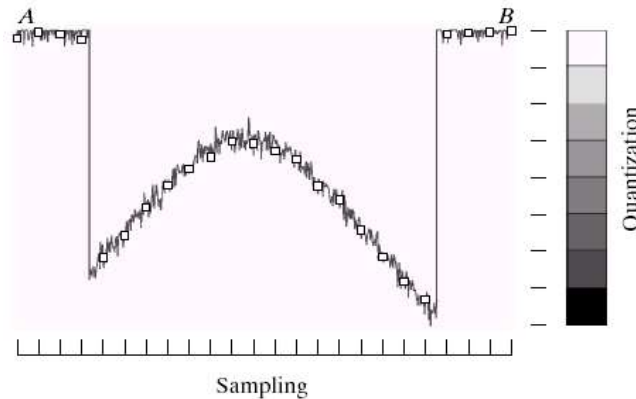
Generating a digital image.

(a) Continuous image.

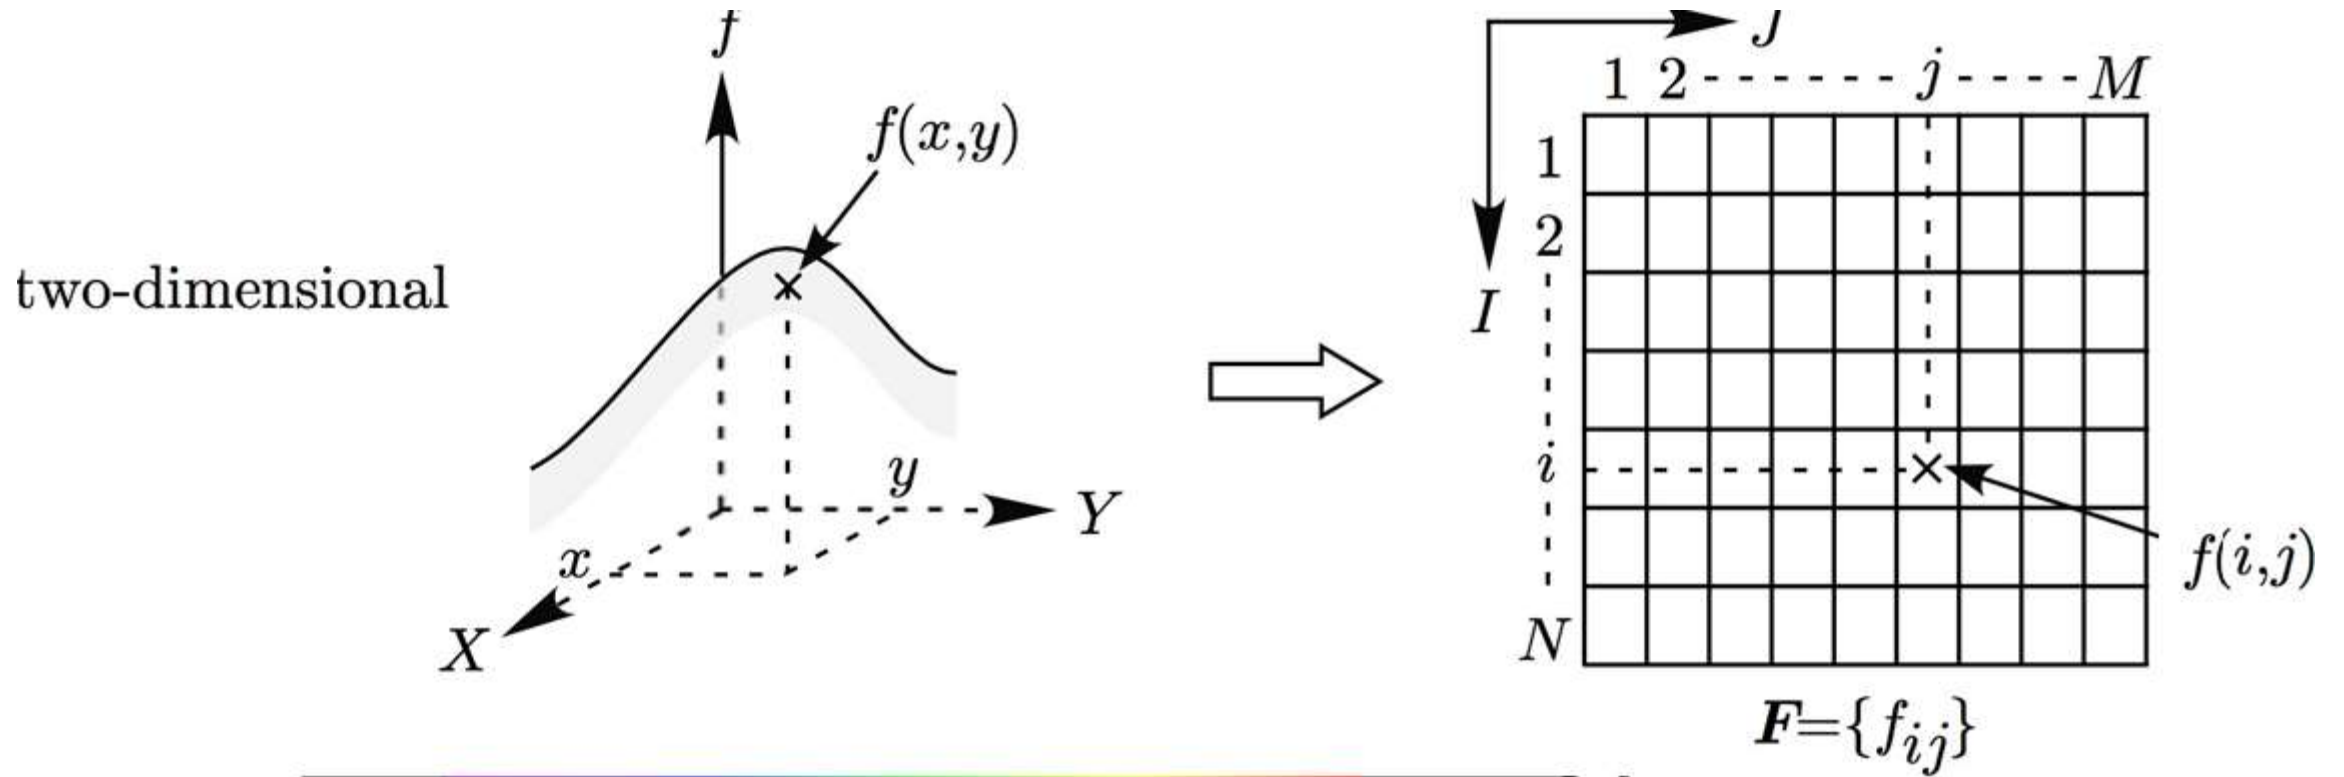
(b) A scaling line from A to B in the continuous image is used to illustrate the concepts of sampling and quantization.

(c) sampling and quantization.

(d) Digital scan line.



Digitization of 2D function



Digitization of 3D function

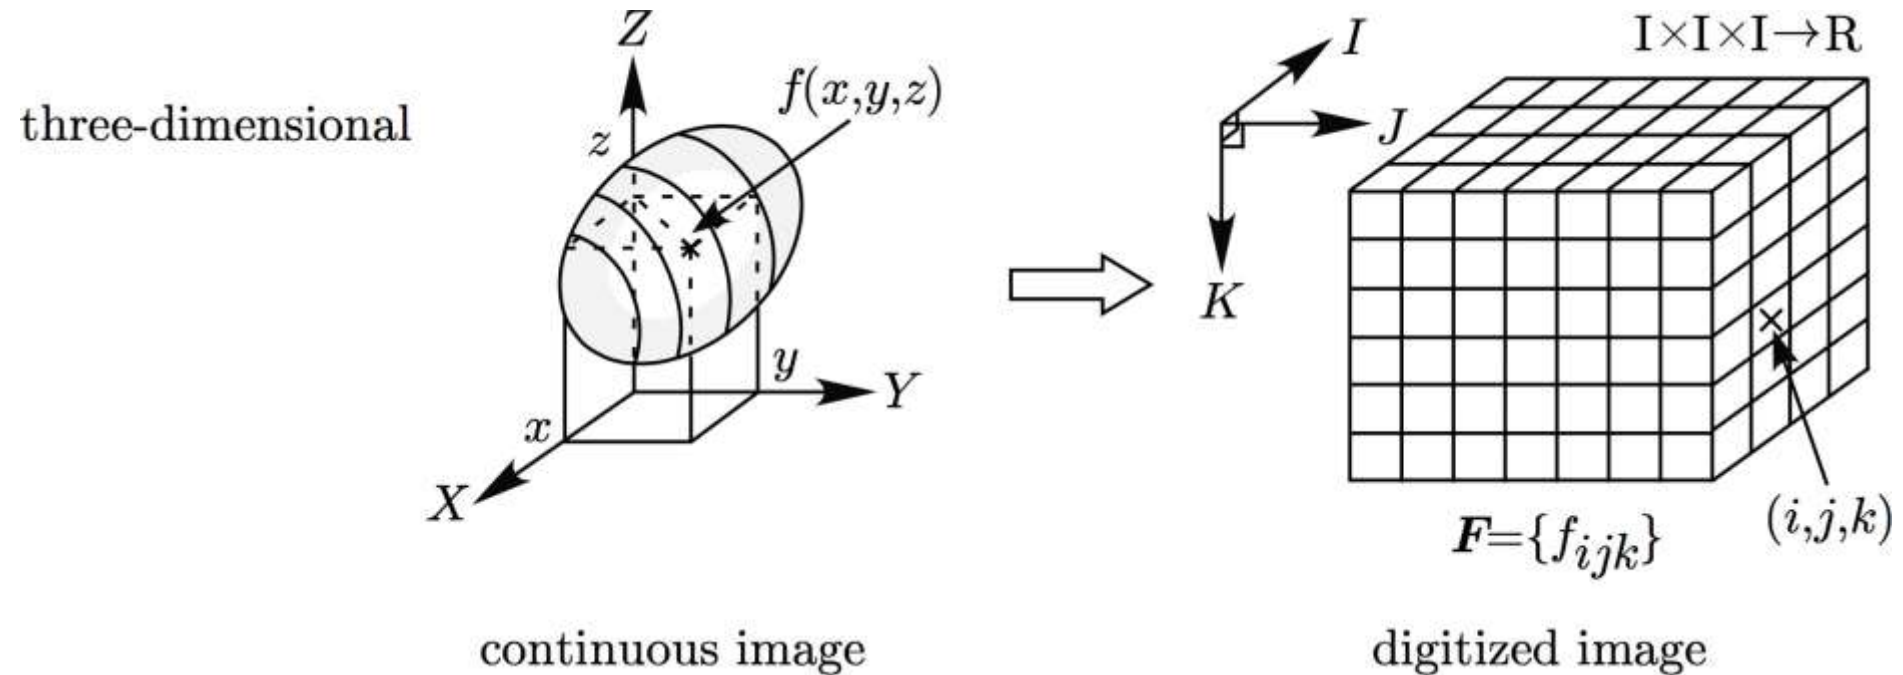


Image as 2D array

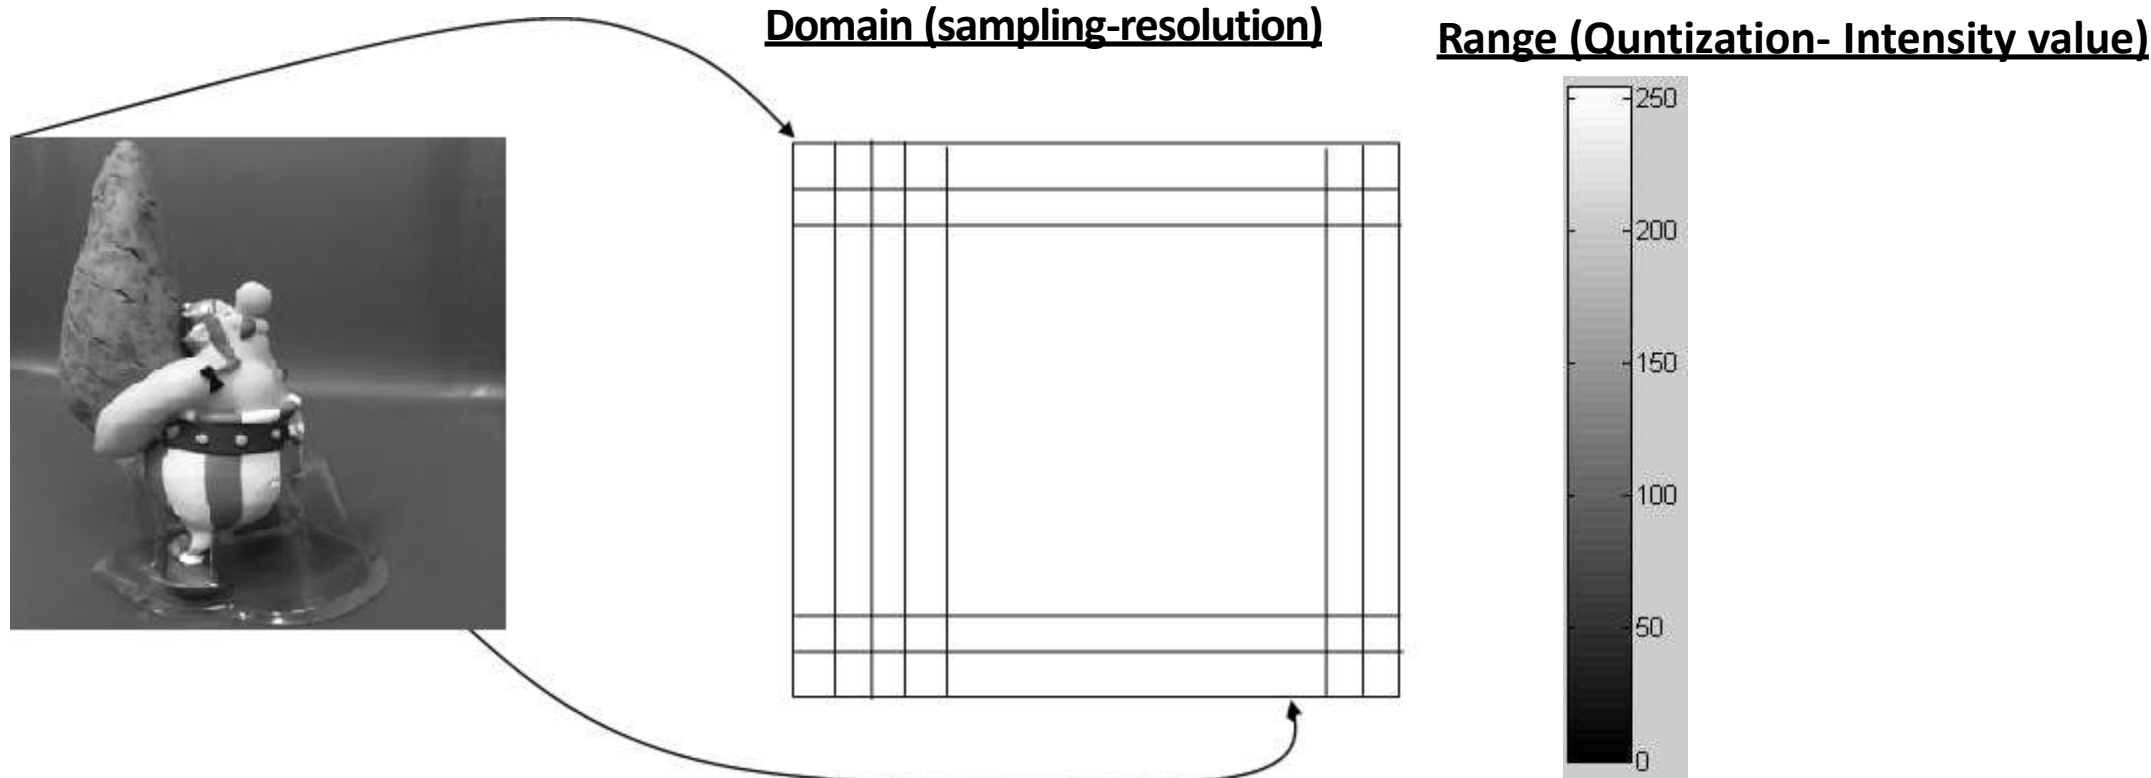


(a) Image

200	202	203	204	205	205	205	206	206	205	205	205	206	203	202	202	202	197	204	198	196	196	199	195	193	198	194	199
202	201	202	204	204	205	205	205	206	206	205	205	205	204	204	204	202	201	199	198	192	198	195	195	196	194	198	
200	201	202	201	204	204	204	205	205	205	205	205	204	202	202	202	204	198	195	195	192	197	193	197	194	195	195	
200	199	201	202	203	203	204	204	205	205	205	205	205	206	204	205	199	203	201	200	200	203	206	195	206	196	193	
199	200	201	201	202	203	203	204	204	205	205	205	204	202	201	201	204	218	199	203	195	200	193	194	195	195	194	
199	200	200	201	202	202	202	203	204	204	204	204	204	205	202	163	140	183	209	194	192	191	197	194	196	193		
199	200	200	201	202	202	202	203	203	204	204	204	204	205	207	282	197	159	99	189	123	180	201	198	197	194	196	
200	200	201	201	202	202	202	203	204	204	204	204	204	201	203	206	210	199	96	107	113	206	199	195	200	196	195	
195	198	200	201	201	202	202	203	202	199	200	207	203	205	204	203	204	201	203	188	198	199	181	173	161	154	191	
198	199	199	200	200	200	201	201	201	202	202	202	202	197	202	196	201	195	161	138	145	11	143	132	106	94		
198	198	198	199	199	199	200	200	200	200	200	200	200	167	83	78	76	74	185	64	84	123	155	133	137	101		
167	137	197	197	198	198	198	198	197	197	194	195	193	168	108	144	156	62	117	138	122	133	57	99	209	191		
195	195	196	196	196	195	196	196	196	195	195	195	195	187	181	174	178	183	125	154	197	190	194	193	198	196		
194	195	195	194	192	191	191	192	193	189	189	197	173	168	166	148	63	58	95	177	172	138	168	188	87	62		
192	192	192	193	192	182	182	183	188	176	169	163	161	154	145	131	29	52	156	172	138	57	79	78	77	75		
185	187	188	182	172	172	172	173	178	163	157	151	141	138	126	97	20	218	158	110	112	21	140	63	50	72		
176	177	176	175	172	162	155	155	145	140	138	122	111	102	114	23	16	235	143	23	135	138	114	108	63			
180	165	158	155	151	141	134	138	133	128	116	105	96	91	95	9	4	217	231	211	95	125	123	119	132	98		
152	150	147	145	141	131	121	113	116	111	110	105	95	87	83	96	120	41	183	209	122	151	41	100	97			
135	133	139	129	125	113	109	1																				

(b) as 2D array

Digitization of image in 2D array



Sampling



1024



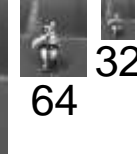
512



256



128



64



32

Resize the sampling images

1024



512



256



128



64



32



Python: `cv2.resize(img, (1024, 1024))`

Quantization



8-bit



7-bit



6-bit



5-bit



4-bit



3-bit



2-bit



1-bit

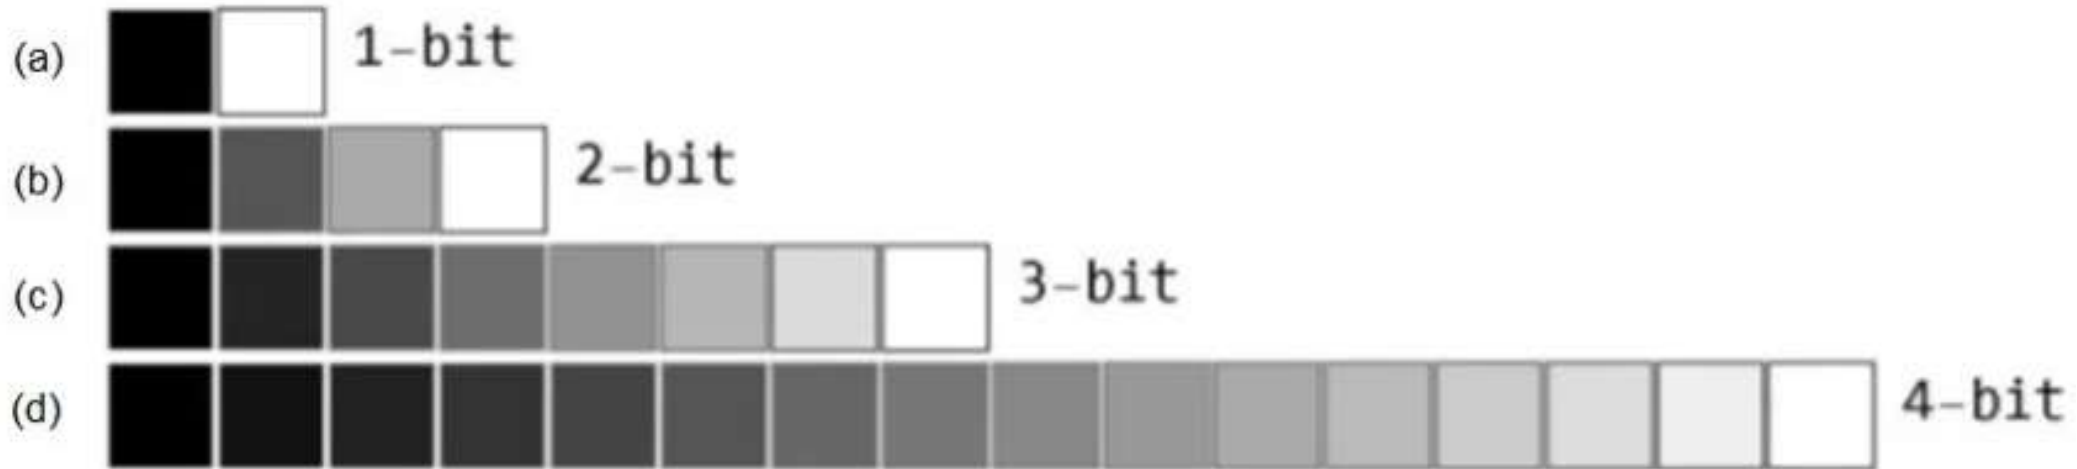
`cv2.kmeans()`

k= 2 for creating 2 color image using the **1-bit level image**.

K= 4 for creating 4 color image using **2-bits level image**

Quantization -Image Intensity

The number of bits used to store that value of the pixel identifies the number of grey levels (or possible grey level values) used to describe a pixel.



Available pixel intensities for, 2-bit, 3-bit, and 4-bit image data

1-bit : $2^1 = 2$ grey level

2-bit : $2^2 = 4$

3-bit : $2^3 = 8$

4-bit : $2^4 = 16$

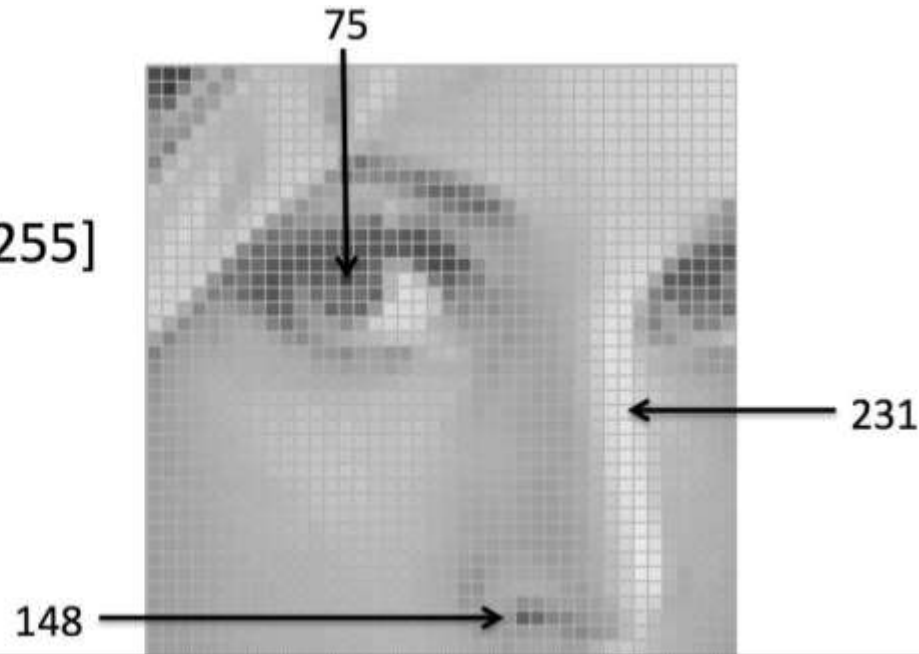
Generally gray images uses 8-bit : $2^8 = 256$

Intensity values after quantization of grayscale image

– Pixel value:

- “grayscale”

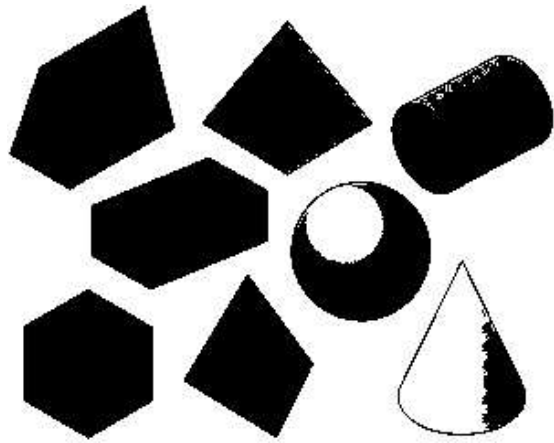
(or “intensity”): [0,255]



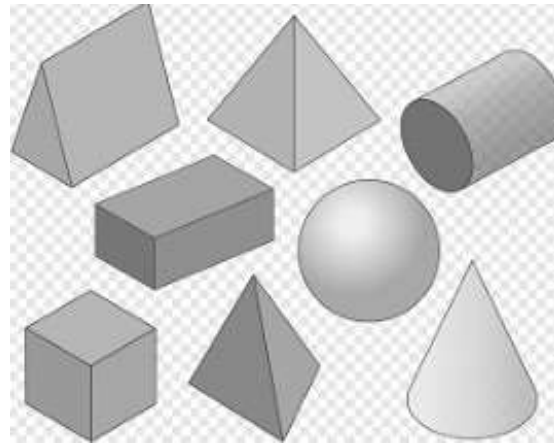
Types of Images

- Binary
- Gray-scale
- Color

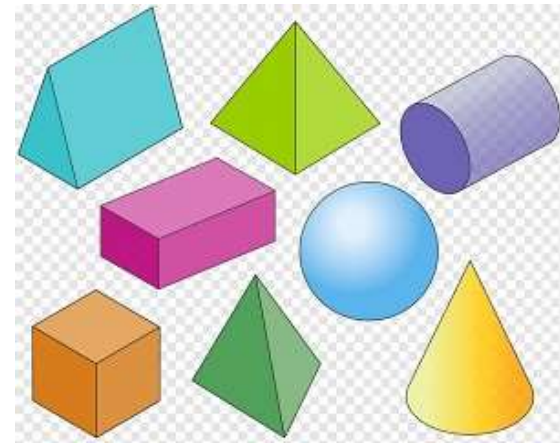
Types of Images



Binary



Gray-scale

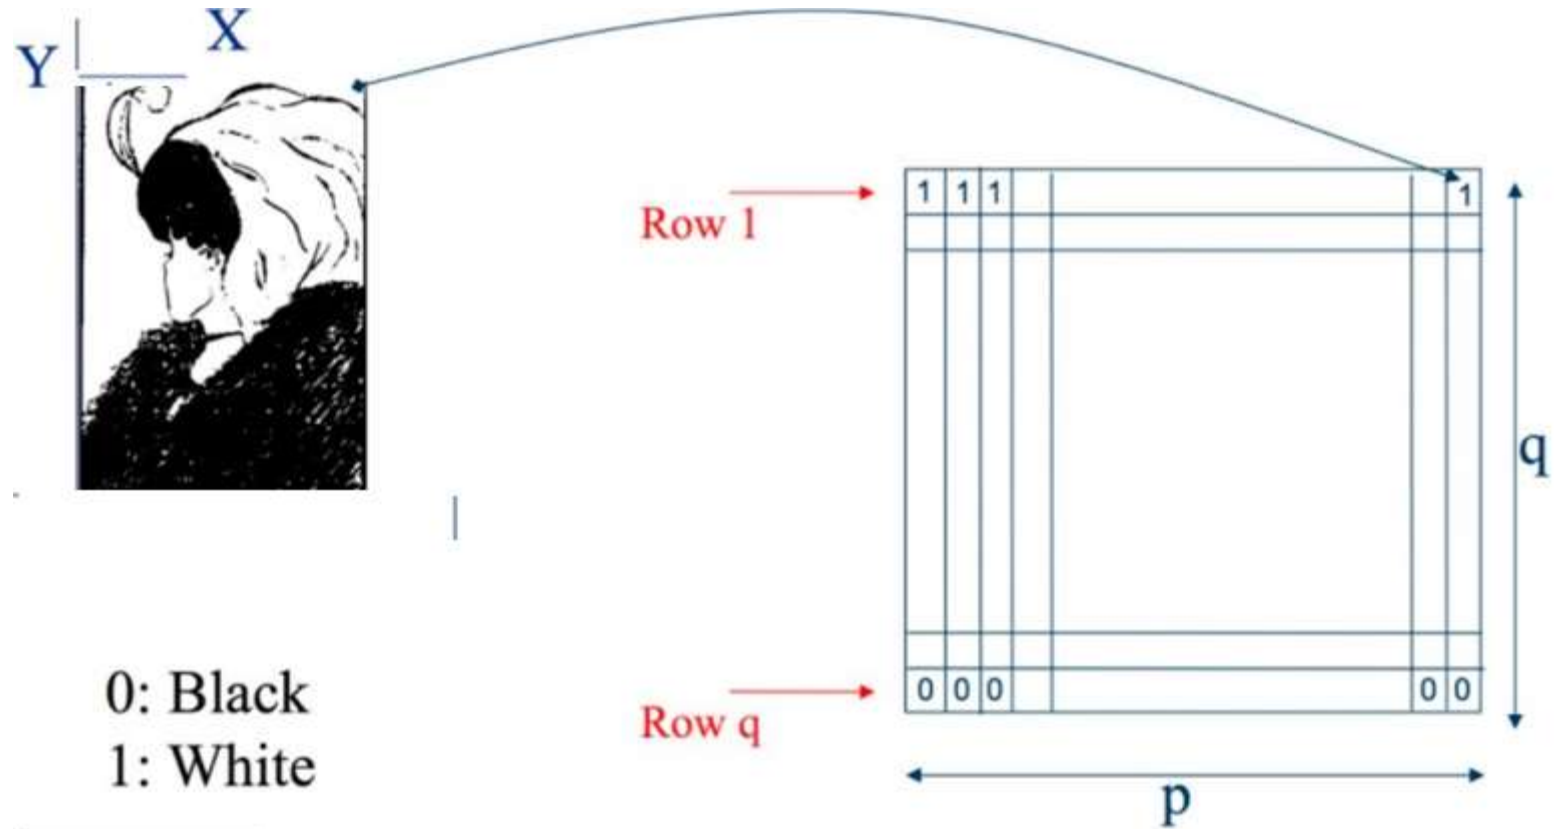


Color

Binary Images

- A binary image is a digital image that has only two possible values for each pixel
- Binary images are also called *bi-level* or *two-level (1bit)*
- Binary images often arise as a result of certain operations such as segmentation, thresholding

Binary Image Representation



Gray-scale Images

- A gray-scale image is composed of shades of gray
- Shades of gray vary from black as the weakest intensity to the white as strongest
- The values of image intensity range from 0 to 255 (8 bits)

Gray-scale Image Representation

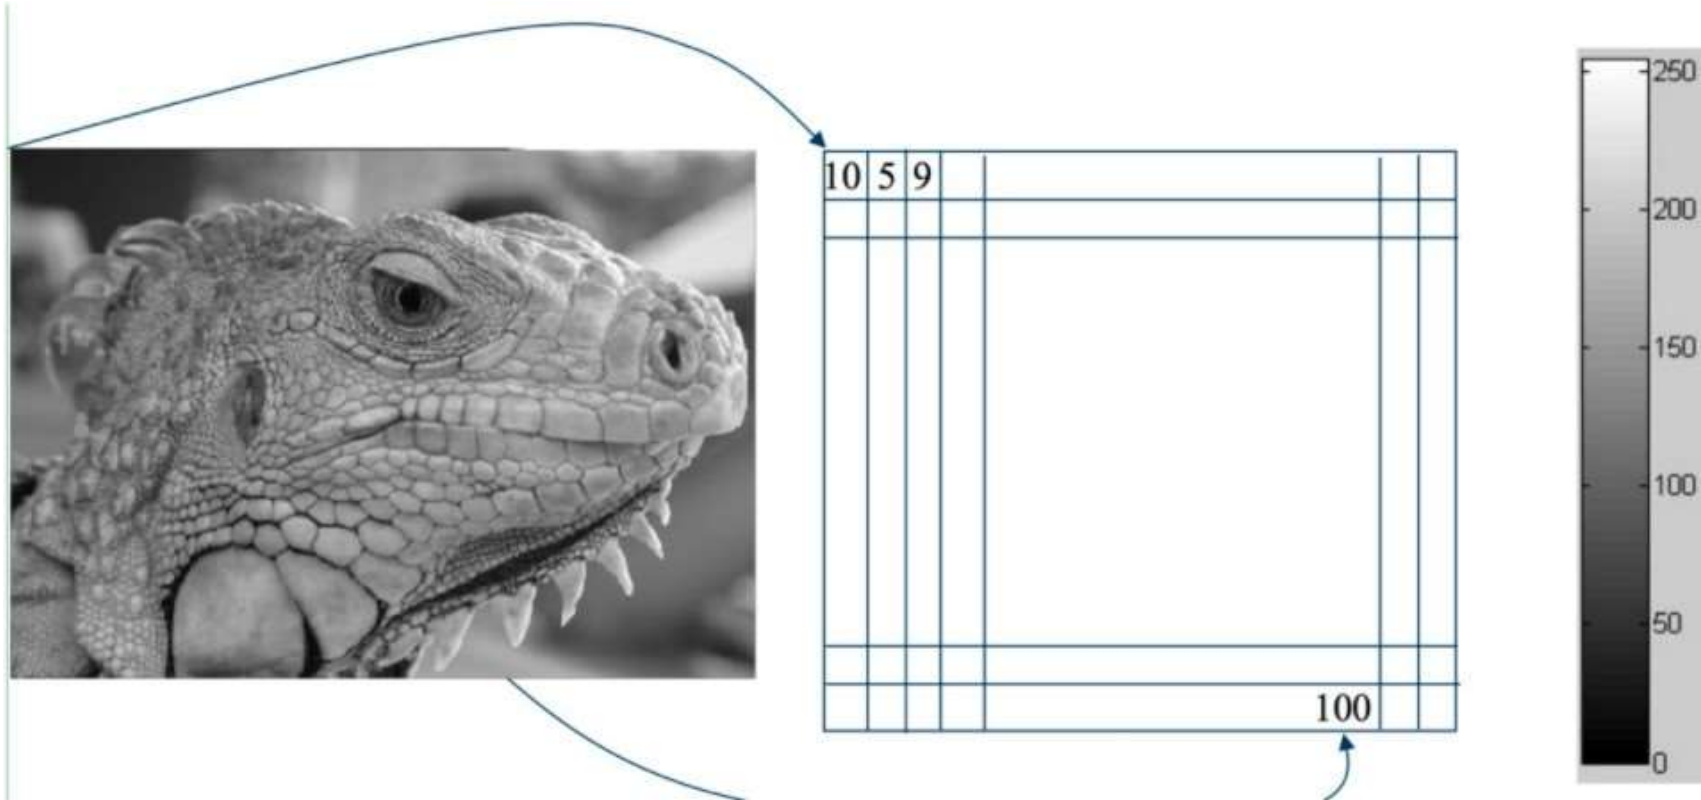


Image credits: Ulas Bagci

Color Images

- Color image is stored as $\text{row} \times \text{cols} \times 3$ array (3 2D array) that defines red, green, and blue color components for each individual pixel
- This model is based on a Cartesian coordinate system

Color Image Representation



B channel

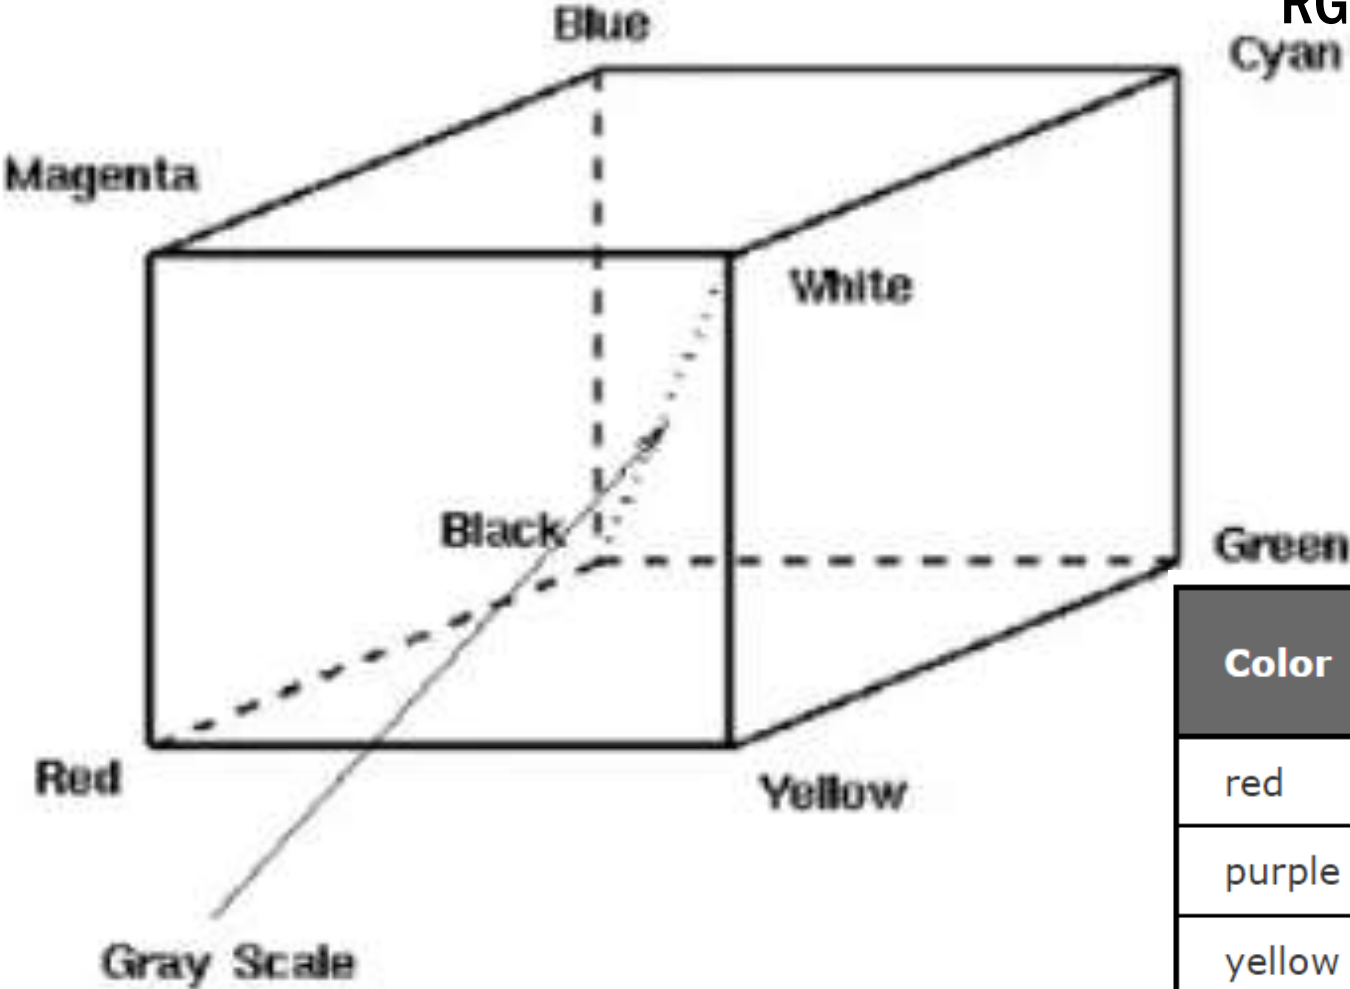


G channel



R channel

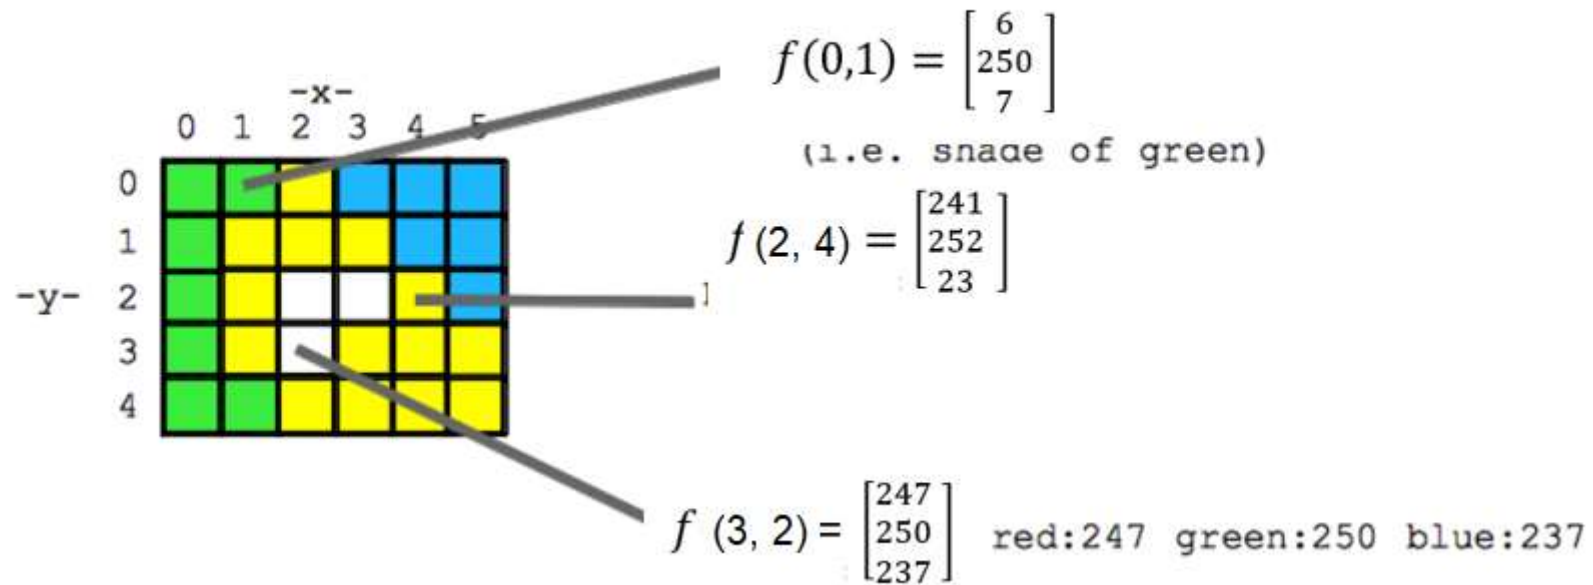
RGB color cube: mixture of RGB color intensities



Color	Red number	Green number	Blue number
red	255	0	0
purple	255	0	255
yellow	255	255	0
dark yellow	100	100	0
white	255	255	255
black	0	0	0

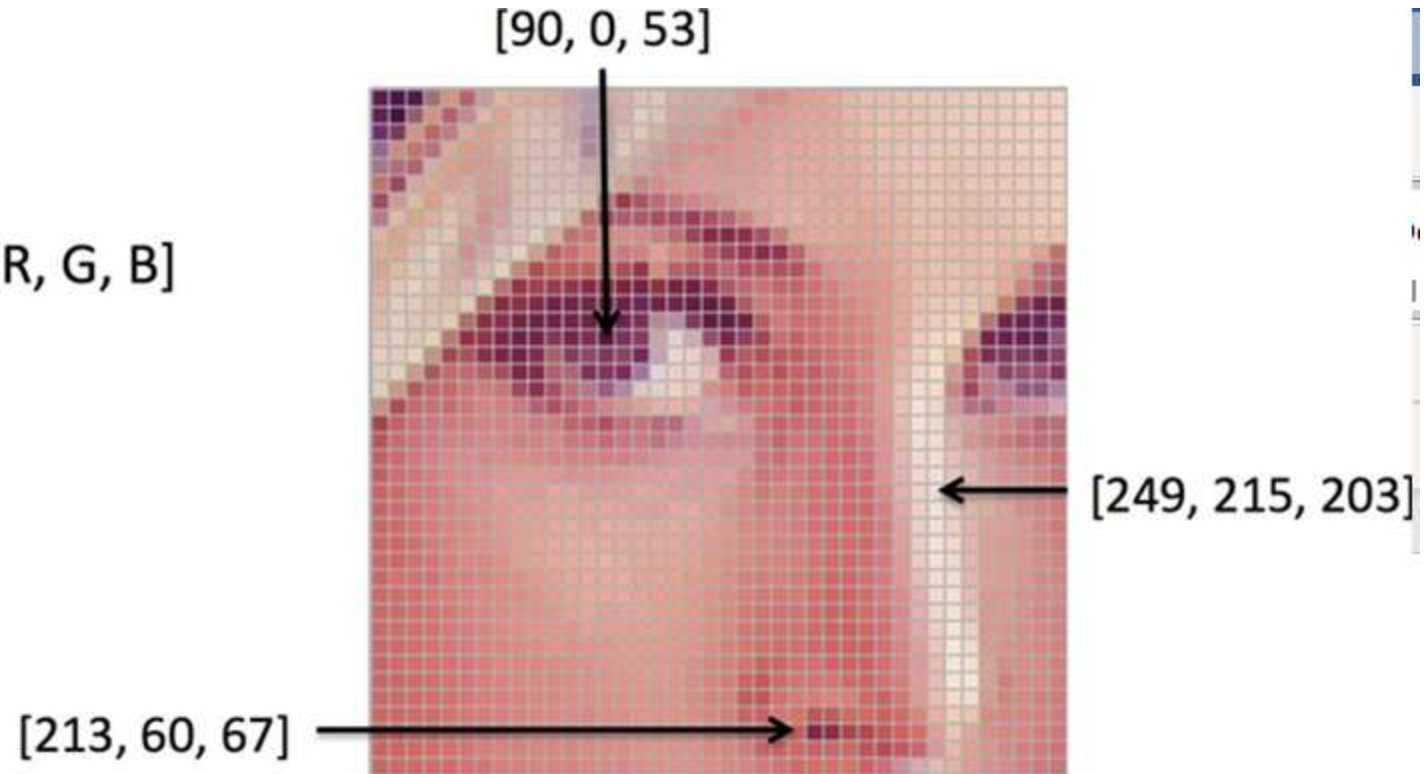
Color image

Pixel values



Color image

- Pixel value:
 - “color”
 - RGB: [R, G, B]



Summary of color intensities

Summary: Color intensities for Gray Scalar and Binary images

- A scalar image has integer values

$$u \in \{0, 1, \dots, 2^a - 1\}$$

- a: level (bit)

- **Ex.** If 8 bit (a=8)

- image spans from 0 to 255
- 0 black and 255 white

- **Ex.** If 1 bit (a=1)

- it is binary image
- 0 and 1 only

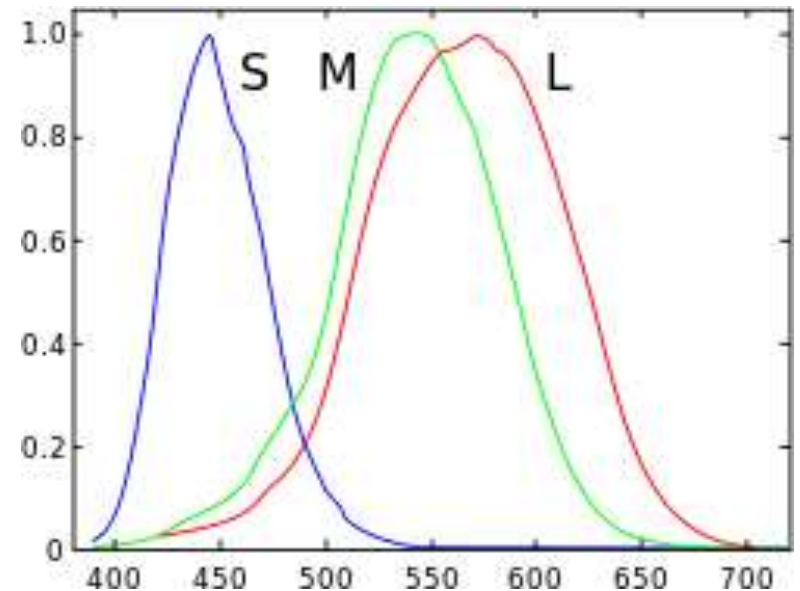


Image Type: RGB (red, green, blue)

- Each channel spans a-bit values.



Human Cone-cells (normalized)
responsivity spectra



What is "Color"?

Human Response to different wavelengths

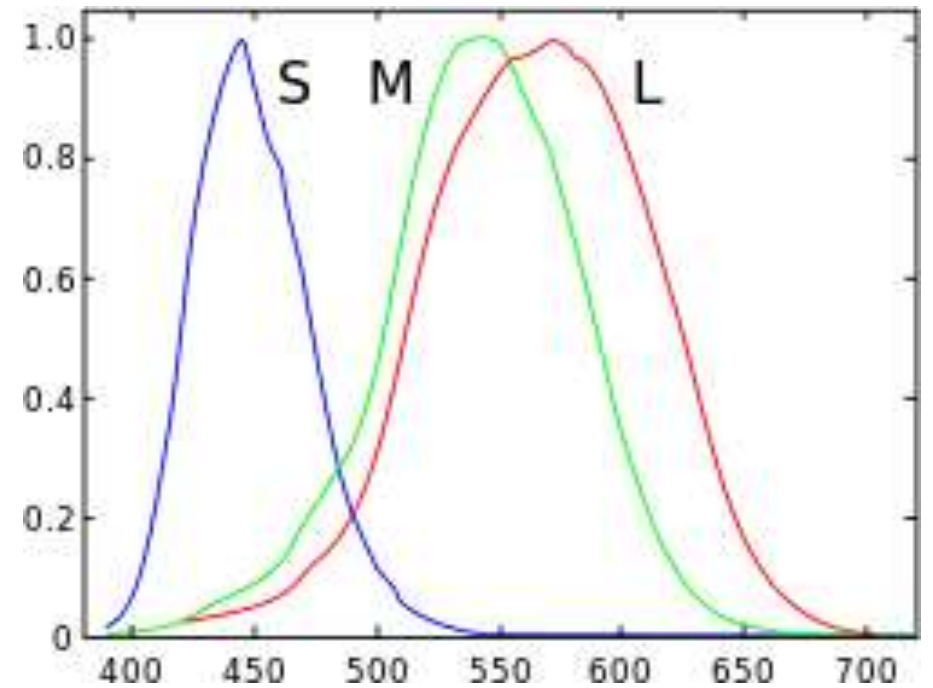
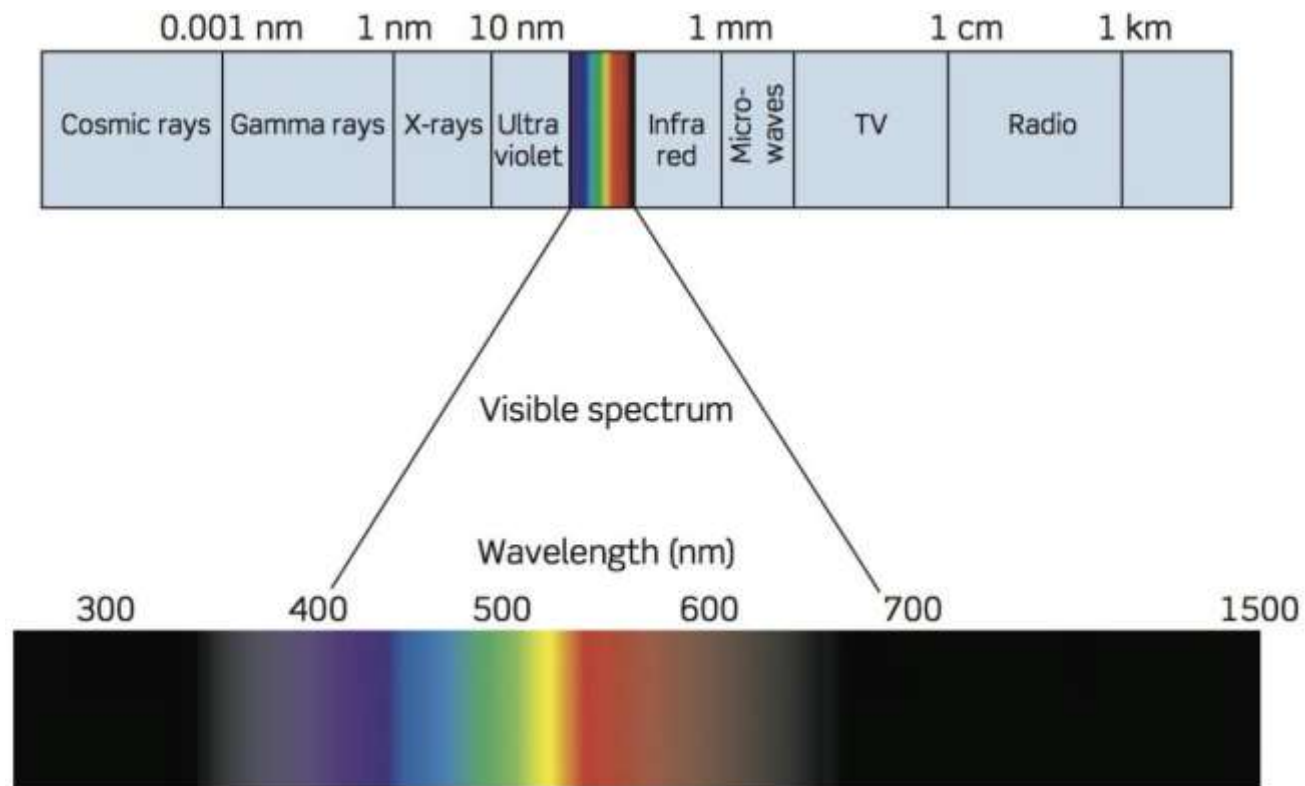
Visible light:



Do We recover spectral distribution $p(\lambda)$?

Sensors in the human eye: Rods & Cones

Color



Color

- If there is **no light**, there is **no color**!
- Human vision can only discriminate a few dozens of grey levels on a screen, but hundreds of thousands of different colors.

- | | |
|-----------------------------|----------------------------|
| • RED -> ~625 to 780 nm | [long wavelength] |
| • ORANGE -> ~ 590 to 625 nm | [long wavelength] |
| • YELLOW -> ~565 to 590 nm | [middle range wavelength] |
| • GREEN -> ~ 500 to 565 nm | [middle range wavelength] |
| • CYAN -> ~485 to 500 nm | [middle range wavelength] |
| • BLUE -> ~440 to 485 nm | [short wavelength] |
| • VIOLET -> ~330 to 440 nm | [very short wavelength] |

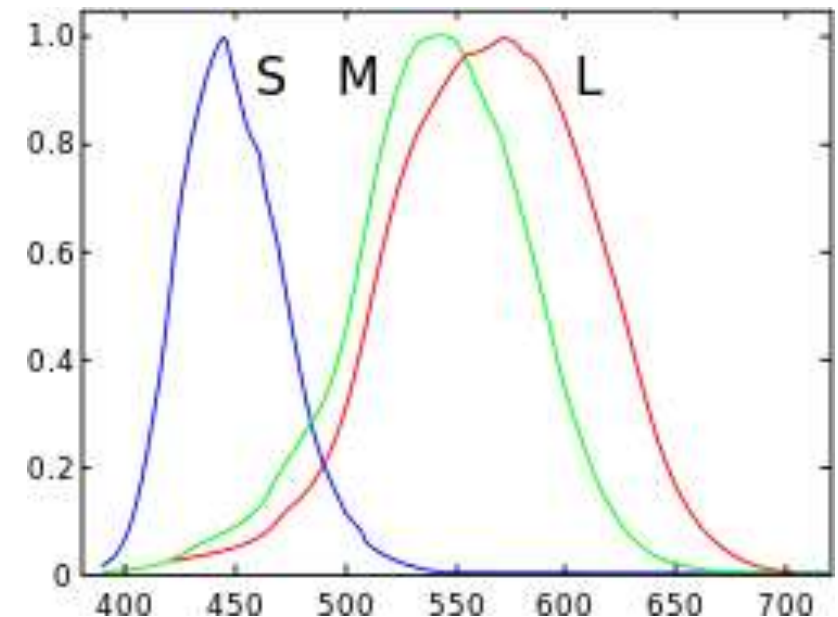
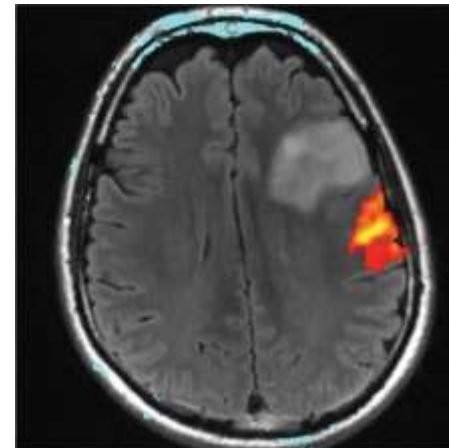
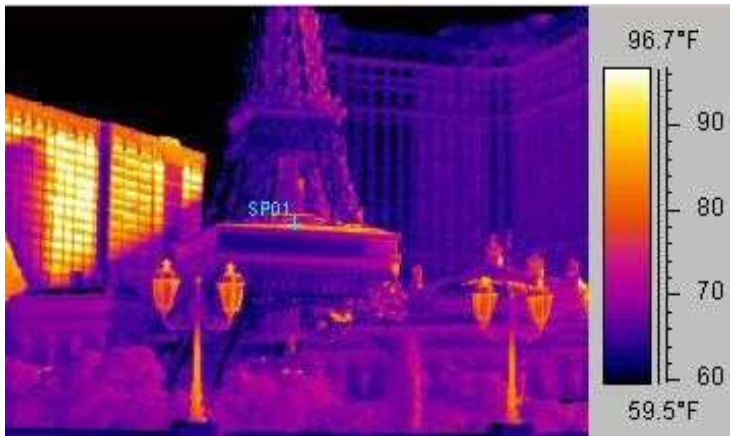


Image (2D matrix) – other examples



Danny Alexander

Image as Function

- $f(x, y)$ gives the intensity at position (x, y)
 - Single intensity information
- A color image is just three functions pasted together
 - $f(x, y) = [f_r(x, y), f_g(x, y), f_b(x, y)]$
 - Vector information

Image processing

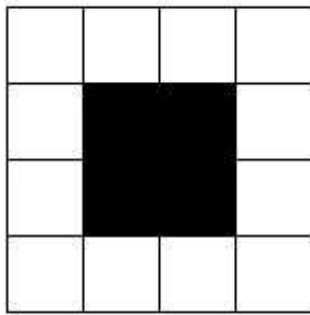
- Image processing usually defines a new image in terms of the existing Image f
- Image processing involves converting an original image into a modified image that is **easier to analyze**.
- Image processing improves images by **fixing problems like noise, motion blur, and defocus blur**.
- $f'(x, y) = g(f(x, y))$ is an example of a simple operation that transforms each image pixel individually

Histograms

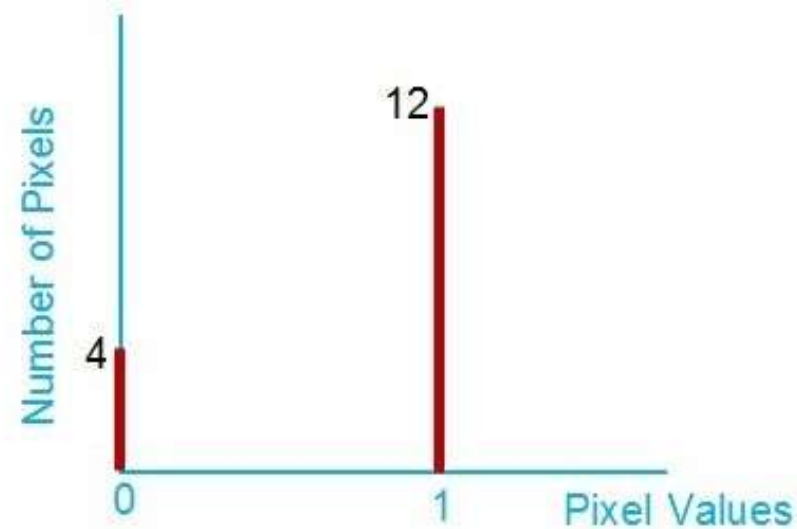
- An image histogram is a **graphical representation of pixel intensities count**
- The x-axis shows the potential intensity values, while the y-axis displays their corresponding counts.

Histogram of a Binary Image

- The first line in the histogram, at gray level 0, shows 4 black pixels in the image.
- The next line indicates 12 white pixels.



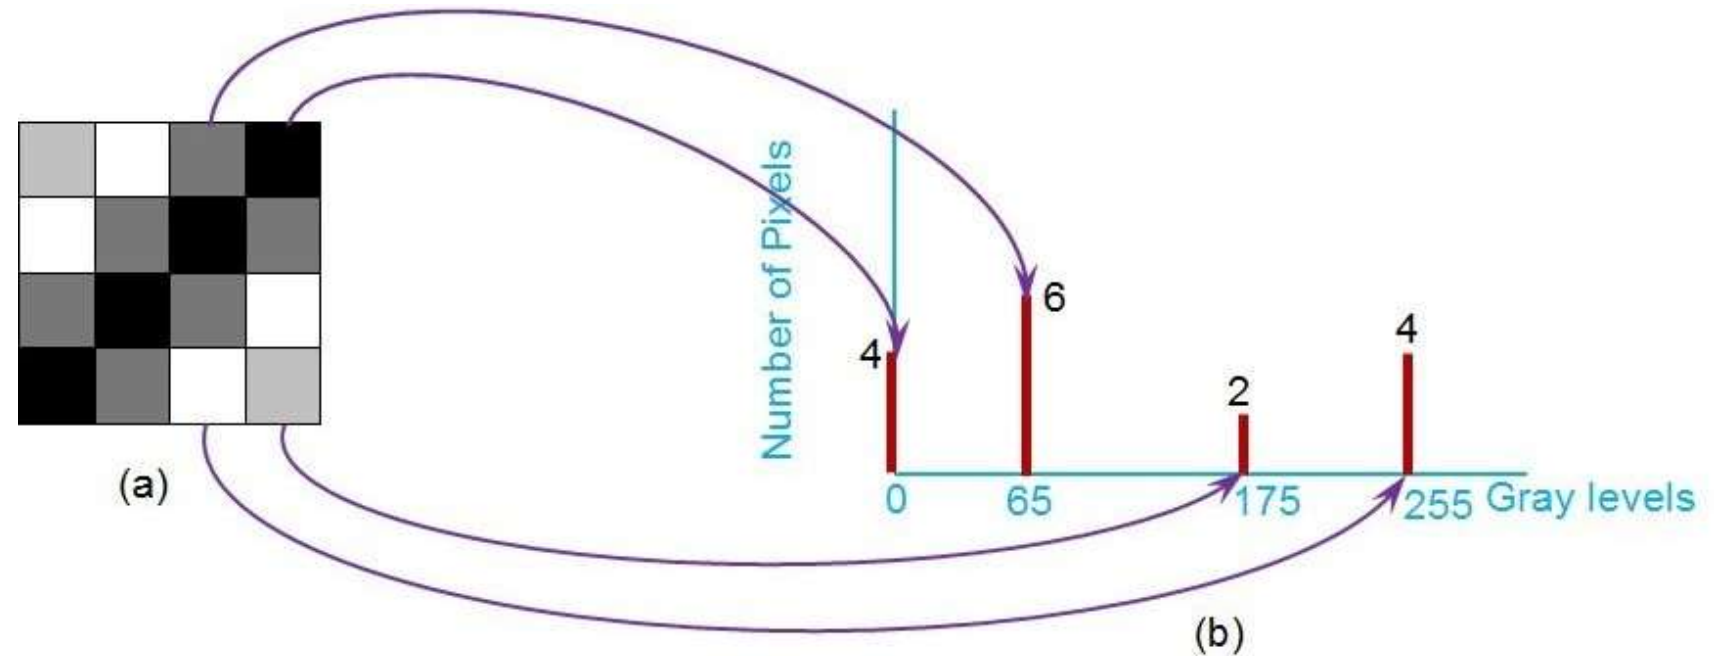
(a)



(b)

Histogram of a Gray-scale Image

- The four-pixel intensities of this image are represented by the four vertical lines of the associated histogram. Pure black, Pure white, dark gray, light gray



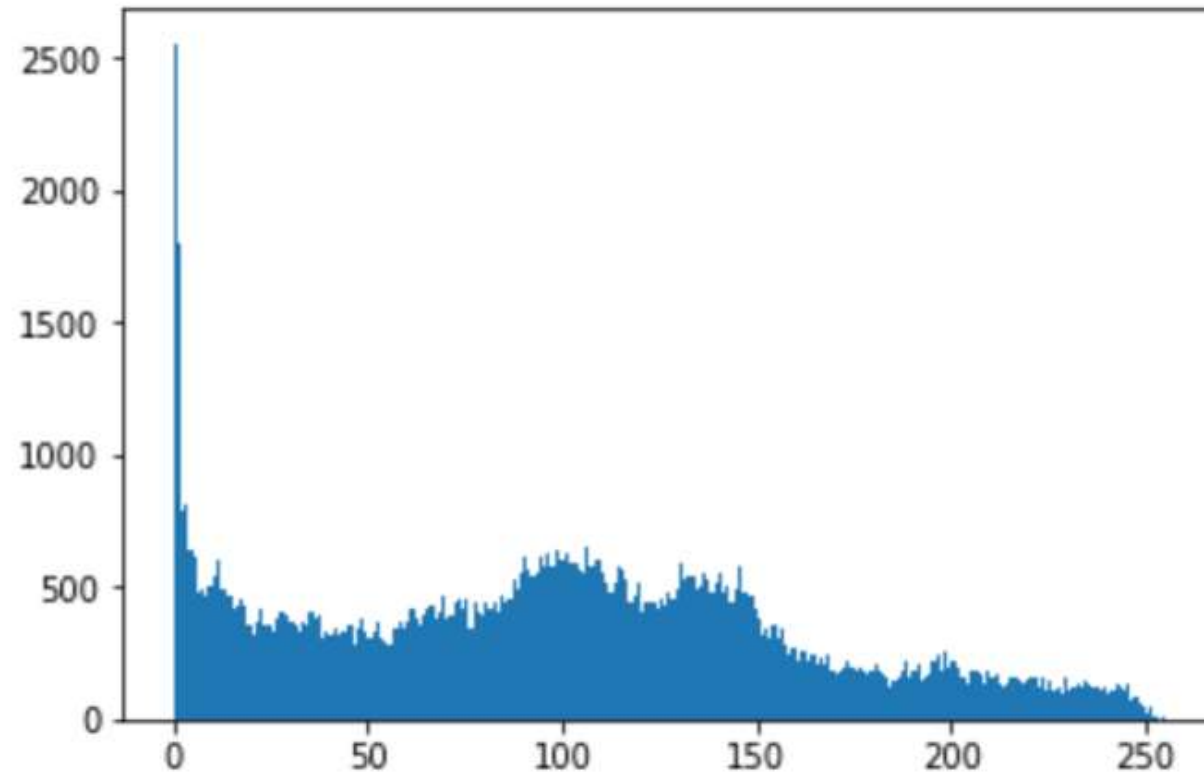
Here the x-axis values span from 0 to 255, which means that there are 256 ($= 2^8$) possible pixel intensities.

Image Histogram- for 8 bit image

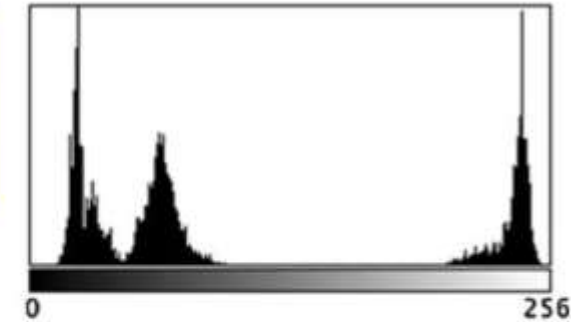
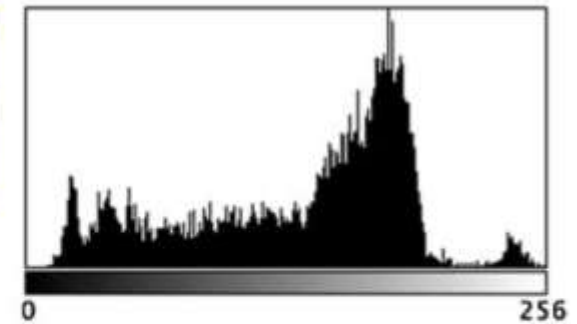
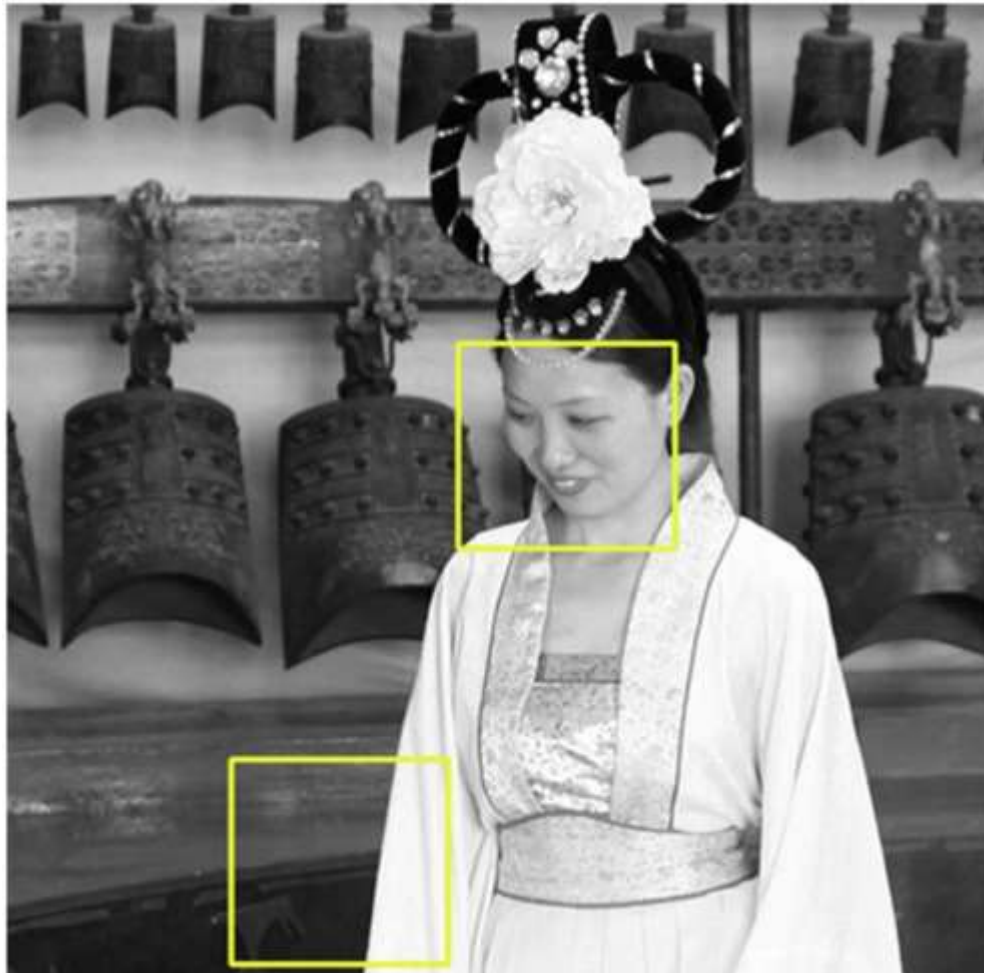


Python code to plot histogram

```
import cv2 as cv
import matplotlib.pyplot as plt
img = cv.imread('parrot.png')
plt.imshow(img)
plt.hist(img.ravel(), bins =255)
plt.show()
```



Histogram Example

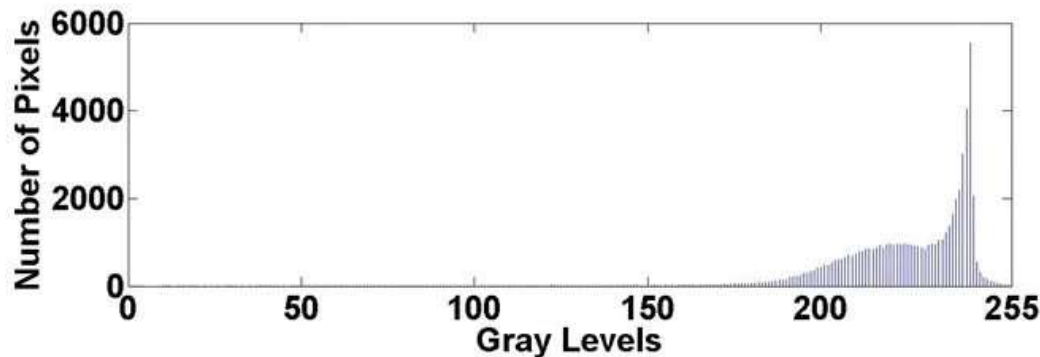


Histogram Applications

- Check image brightness
- Check image darkness
- Histogram equalization to enhance the image
- Apply thresholding – generate binary image

Check Image Brightness

- One can get a general idea of the brightness of an image by looking at the histogram
- If the histogram values are concentrated toward the right, the image is brighter



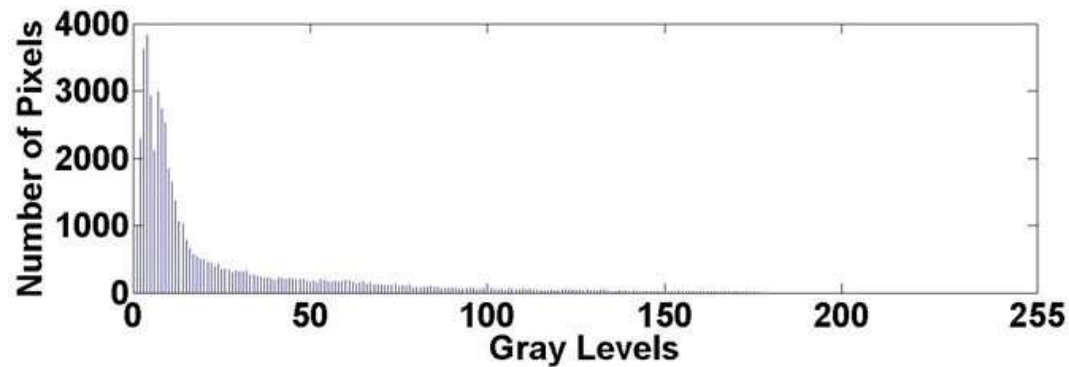
(a)



(b)

Check Image darkness

- If the histogram values are concentrated toward the left, the image is darker



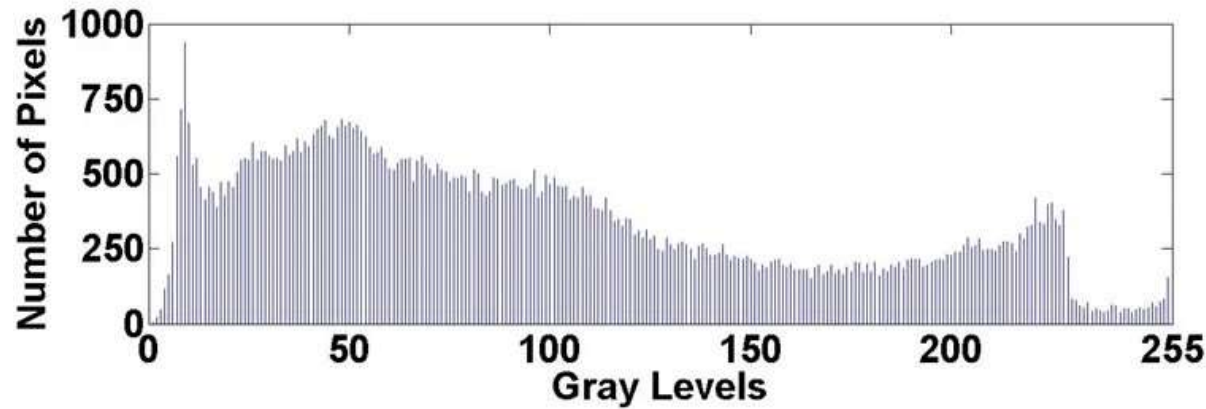
(a)



(b)

The contrast of the Image

A histogram in which the **pixel counts** evenly cover a **broad range** of **gray-scale levels** indicates an image with good/high contrast



(a)



(b)

Figure 1: High contrast image

The contrast of the Image

A histogram in which the **pixel counts are not evenly cover a broad range of gray-scale levels** indicates an image with *low contrast*

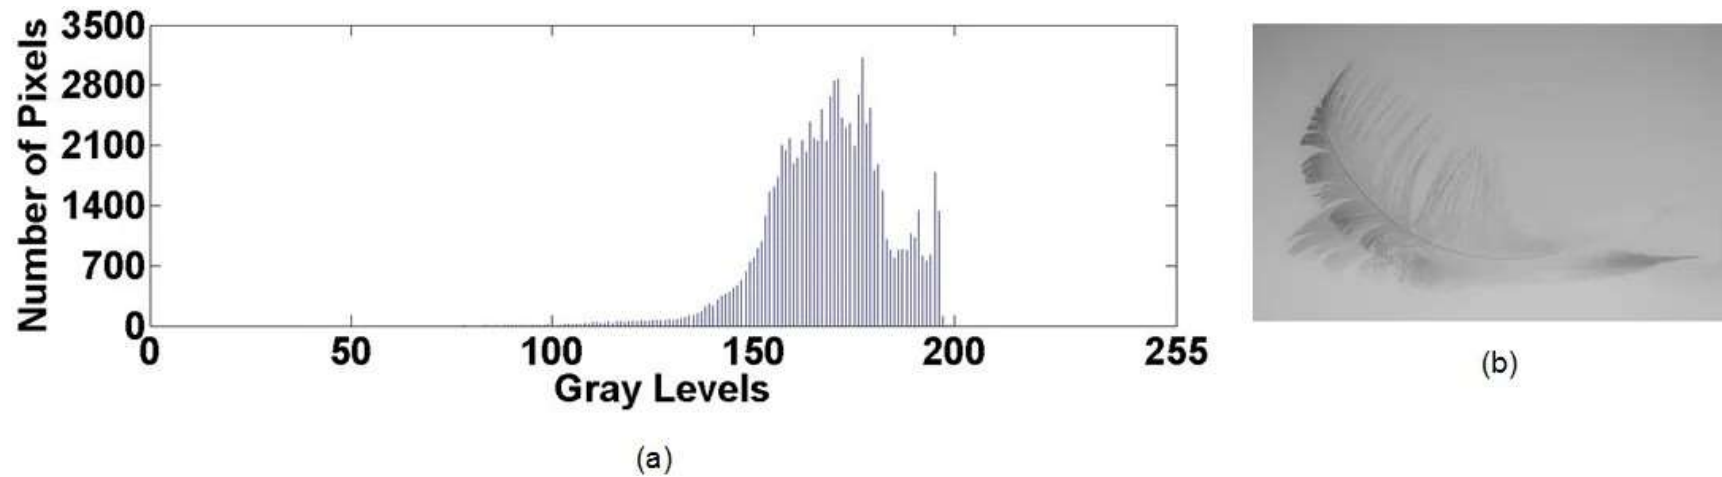


Figure 2: Low contrast image

Histogram Equalization for Image Enhancement

- Enhance the quality of the image by spreading the most frequent intensity values using the following steps:
 1. Calculate the frequency against each intensity value
 2. Compute the probability of each frequency
 3. Compute Cumulative Probability or Cumulative Distribution Function (CDF)
 4. Multiply the CDF with the required intensity levels
 5. Round the resultant value to floor

Histogram Equalization for Image Enhancement

Suppose this matrix is representing pixel intensities

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

- Intensity of pixels range between 1 – 8
- Suppose we want to apply histogram equalization and scale the intensity to 1 - 20

Histogram Equalization for Image Enhancement

Example

1. Calculate the frequency against each intensity value

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

Intensity	1	2	3	4	5	6	7	8
Count	1	3	3	2	2	1	3	1

Histogram Equalization for Image Enhancement

Example

2. Compute the probability of each frequency

$$probability = \frac{\text{No. of pixels associated to a particular intensity}}{\text{Total pixels}}$$

Histogram Equalization for Image Enhancement

Example

2. Compute the probability of each frequency

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

Intensity	1	2	3	4	5	6	7	8
Count	1	3	3	2	2	1	3	1
Probability	.062	.187	.187	.125	.125	.062	.187	.062

Histogram Equalization for Image Enhancement

Example

3. Compute Cumulative Distribution Function (CDF) i.e. cumulative probabilities

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

Intensity	1	2	3	4	5	6	7	8
Count	1	3	3	2	2	1	3	1
Probability	.0625	.1875	.1875	.125	.125	.0625	.1875	.0625
Cumulative Probability	.0625	.25	.4375	.5625	.6875	.75	.9375	1

Histogram Equalization for Image Enhancement

Example
4. Multiply the CDF with the required intensity levels

- ▶ As we want to increase the intensity to the range 1-20, multiply each cumulative probability by 20

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

Intensity	1	2	3	4	5	6	7	8	
Count	1	3	3	2	2	1	3	1	
Probability		.0625	.1875	.1875	.125	.125	.0625	.1875	.0625
Cumulative Probability		.0625	.25	.4375	.5625	.6875	.75	.9375	1
CP x 20		1.25	5	8.75	11.25	13.75	15	18.75	20

Histogram Equalization for Image Enhancement

Example

5. Round the resultant value to the floor

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

Intensity	1	2	3	4	5	6	7	8
Count	1	3	3	2	2	1	3	1
Probability	.0625	.1875	.1875	.125	.125	.0625	.1875	.0625
Cumulative Probability	.0625	.25	.4375	.5625	.6875	.75	.9375	1
CP x 20	1.25	5	8.75	11.25	13.75	15	18.75	20
Floor Rounding	1	5	8	11	13	15	18	20

Histogram Equalization for Image Enhancement

Example

Original Image has been transformed to equalized image with different intensities

3	2	4	5
7	7	8	2
3	1	2	3
5	4	6	7

Before

8	5	11	13
18	18	20	5
8	1	5	8
13	11	15	18

After

Python function for Histogram Equalization

```
# Importing libraries
```

```
import cv2 as cv
```

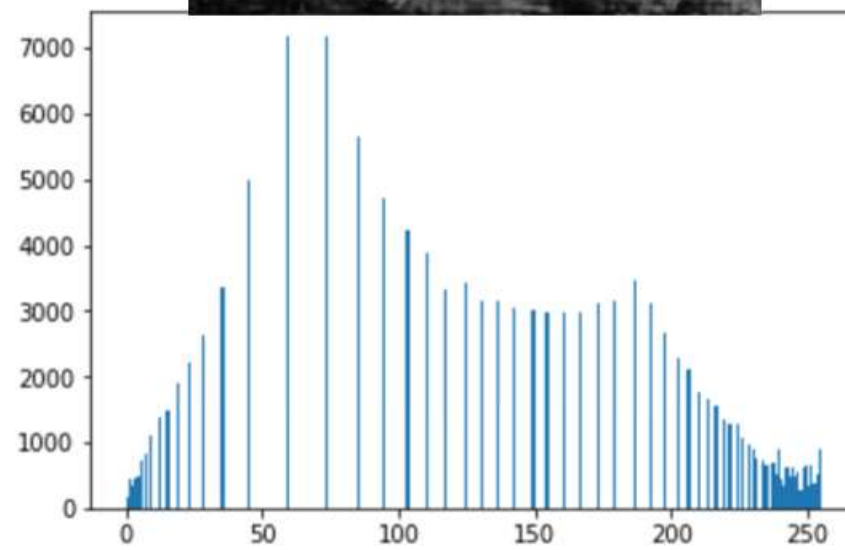
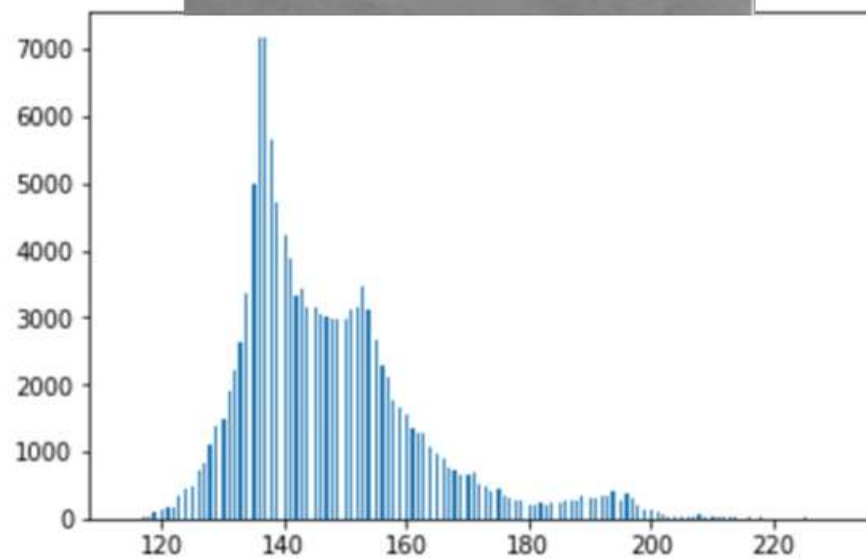
```
# OpenCV function
```

```
cv.equalizeHist(image)
```

```
# Importing libraries
import cv2
import matplotlib.pyplot as plt
def histogram_equalize(img):
    b, g, r = cv2.split(img)
    red = cv2.equalizeHist(r)
    green = cv2.equalizeHist(g)
    blue = cv2.equalizeHist(b)
    return cv2.merge((blue, green, red))

img = cv2.imread('eq.jpg')
plt.imshow(img)
img.shape

equ = histogram_equalize(img)
equ.shape
res = np.hstack((img,equ)) #stacking images side-by-side
plt.imshow(res)
```



Take-Home Quiz 1

a. Read and convert color(RGB) image to gray-scale image

$$f'(x, y) = g(f(x, y))$$

where $f'(x, y)$ is the gray-scale image and f in representing a color image, whereas g is a function that is used to convert the color image to gray-scale

b. Apply the provide code (shown in previous slides) to enhance contrast of any low contrast image (dark or bright)

Intensity profiles for selected (two) rows

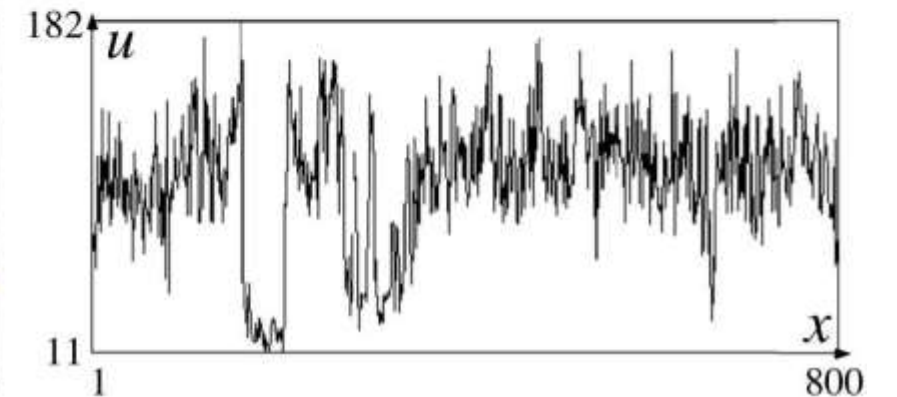
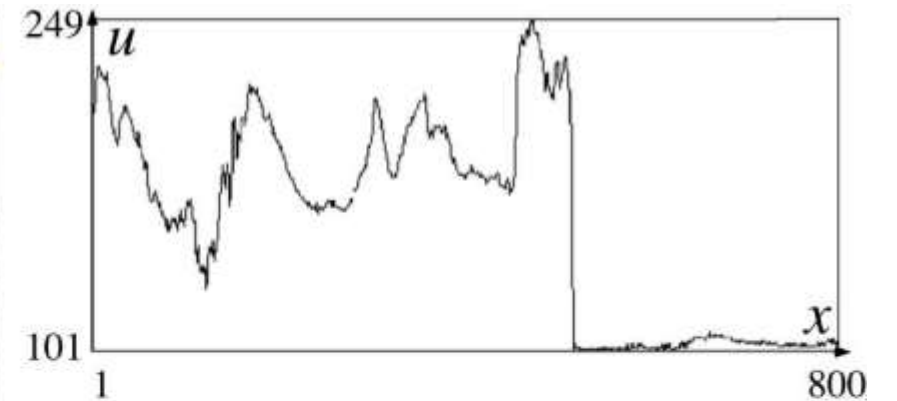


Image noise

- Light Variations
 - Camera Electronics
 - Surface Reflectance
 - Lens
-
- Noise is random,
 - it occurs with some probability
 - It has a distribution

Noise

- $I_{\text{original}}(x,y)$ – true pixel value at (x,y)
- $n(x,y)$ - noise at (x,y)
- $I_{\text{observed}}(x,y) = I_{\text{original}}(x,y) + n(x,y)$ additive noise



Noise

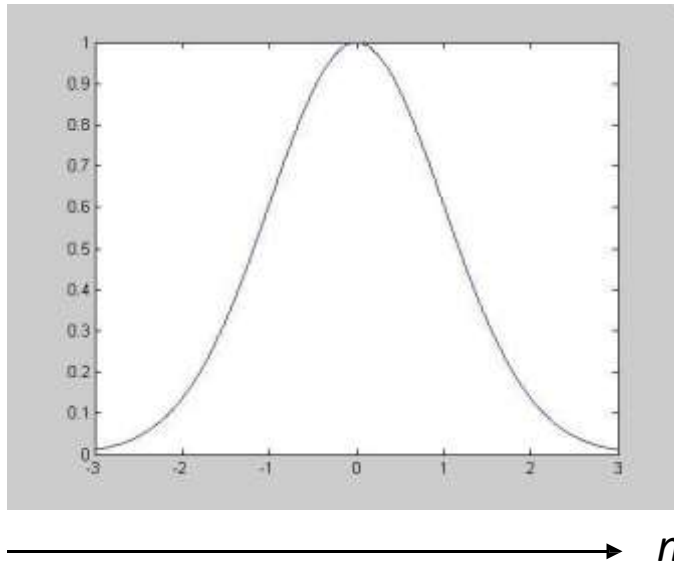
- $I_{\text{original}}(x,y)$ – true pixel value at (x,y)
- $n(x,y)$ - noise at (x,y)
- $I_{\text{observed}}(x,y) = I_{\text{original}}(x,y) * n(x,y)$ multiplicative noise



Gaussian Noise

$$n(x, y) \approx g(n) = e^{\frac{-n^2}{2\sigma^2}}$$

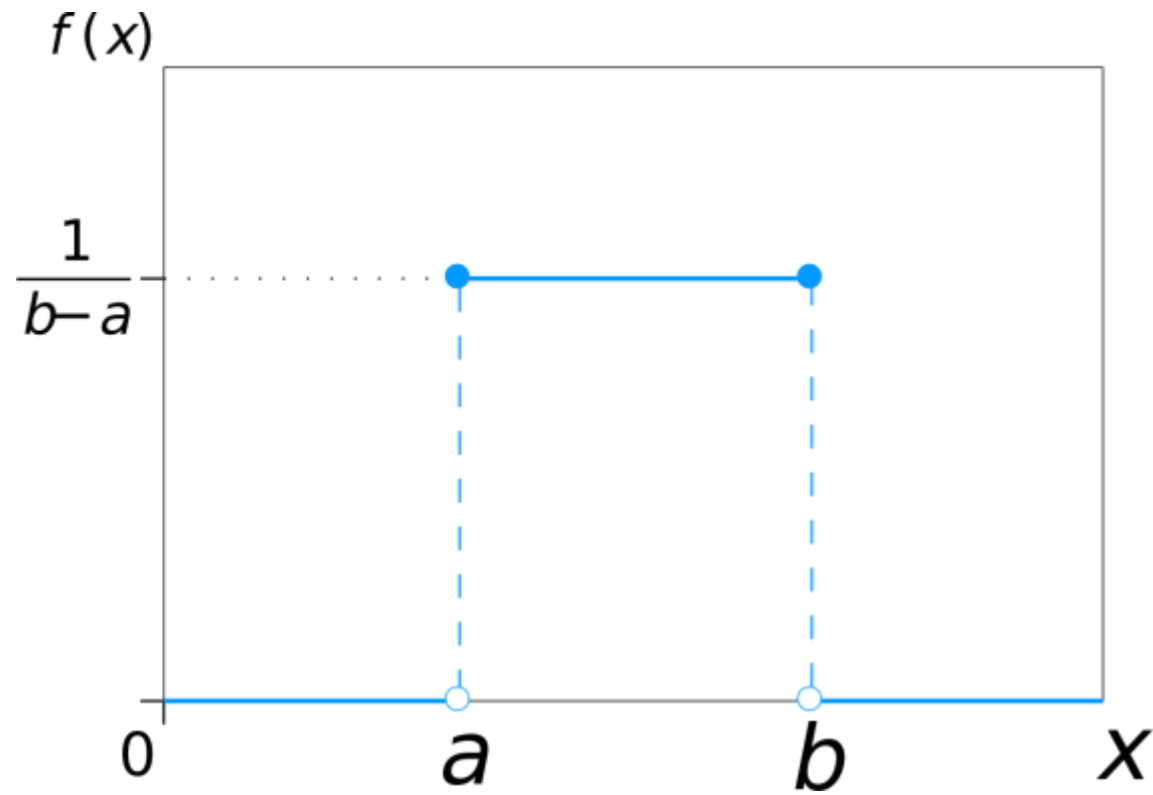
$g(n)$



Probability Distribution
 n is a random variable



Uniform distribution



Salt and pepper noise

- Each pixel is randomly made black or white with a uniform probability distribution



Salt-pepper

So.... How do we detect an
object in an
image?

Naïve approach: Template Matching

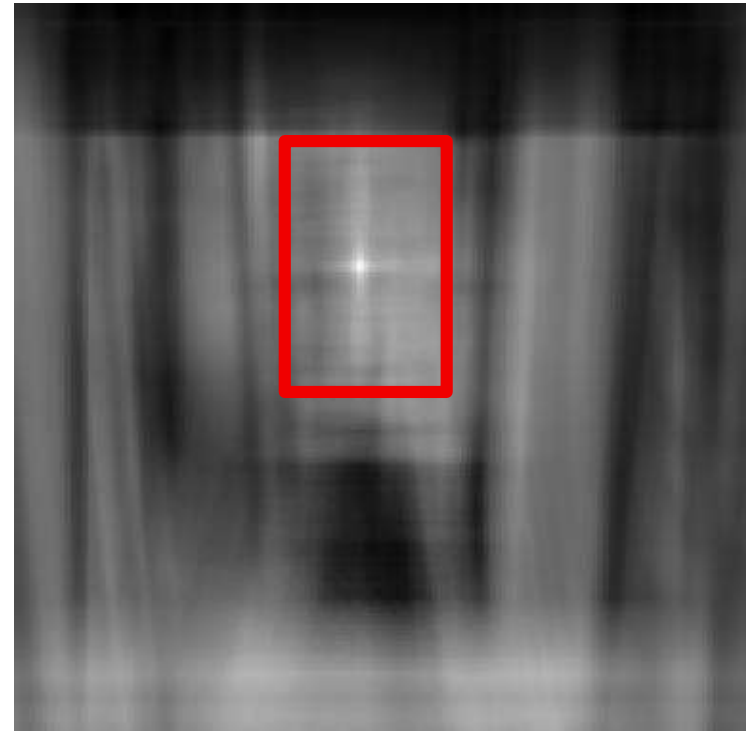
This is a chair



Find the chair in this image



Output of correlation



Template Matching

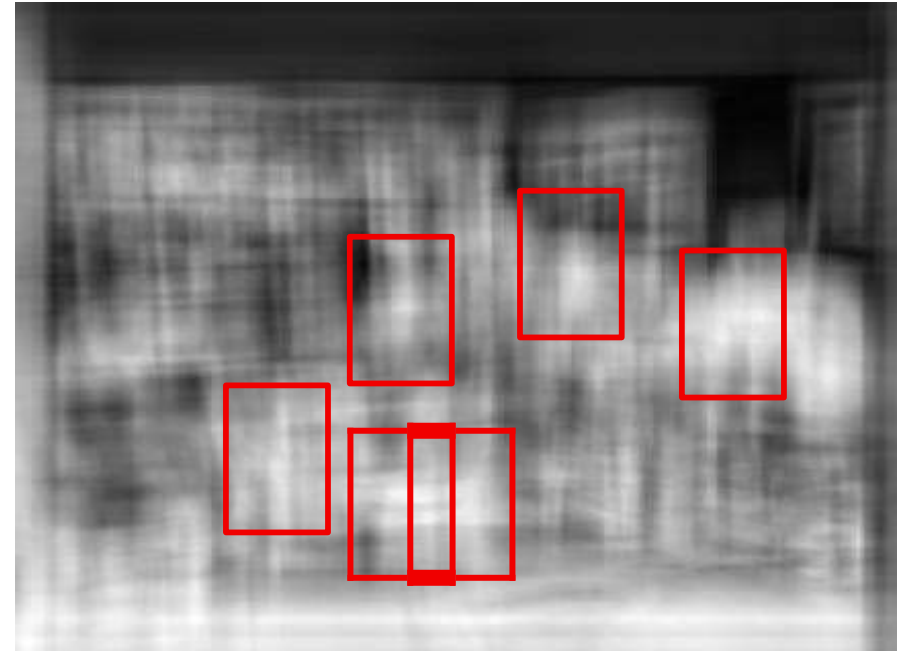
Find the chair in this image



Template Matching



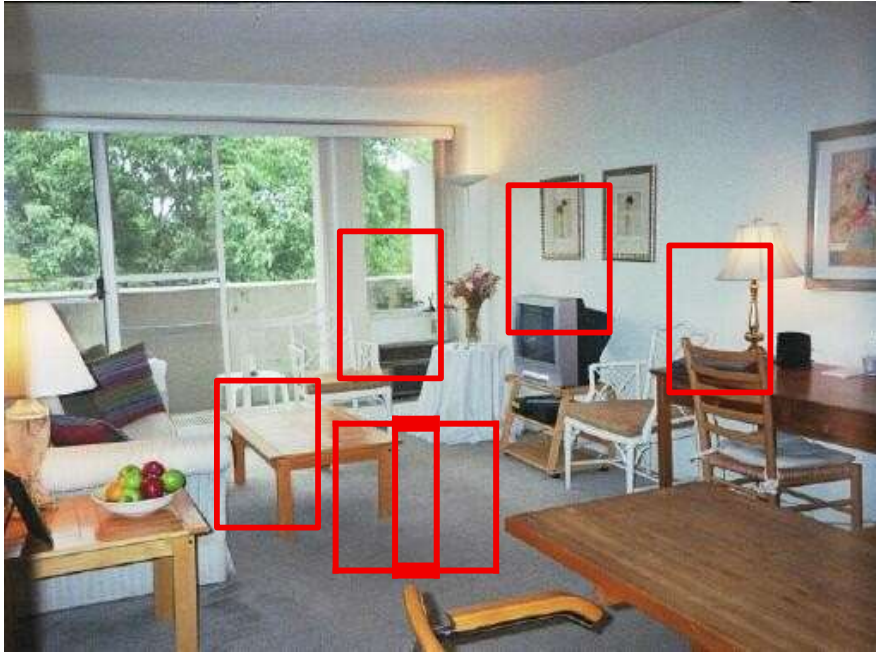
Find the chair in this image



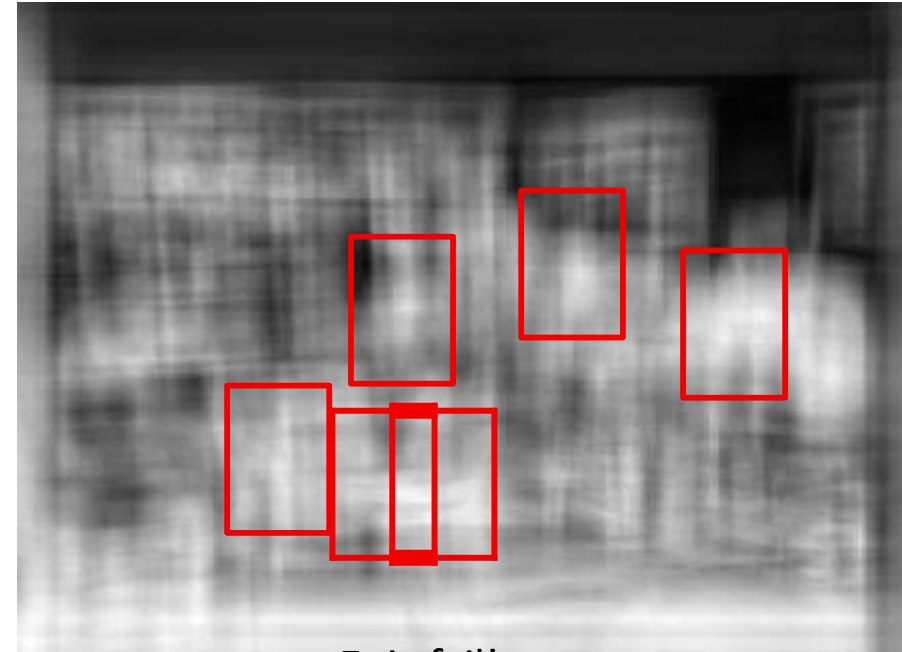
Epic fail!

Simple template matching is not going to make it

Template Matching



Find the chair in this image



Epic fail!

Simple template matching is not going to make it

Idea

- Instead of comparing raw image pixels:
 - First map those pixels into another (more robust) form,
 - And then compare those mapped forms.
 - Finally, select the closest image map (how do you define “closest”? Metrics).
- Features, examples:
 - compute edges
 - compute color histograms
 - gradients
 - HOG
 - SIFT
 - ...

Image filtering

- Image filtering: compute function of local neighborhood at each position

h=output f=filter I=image

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

2d coords=k,l 2d coords=m,n

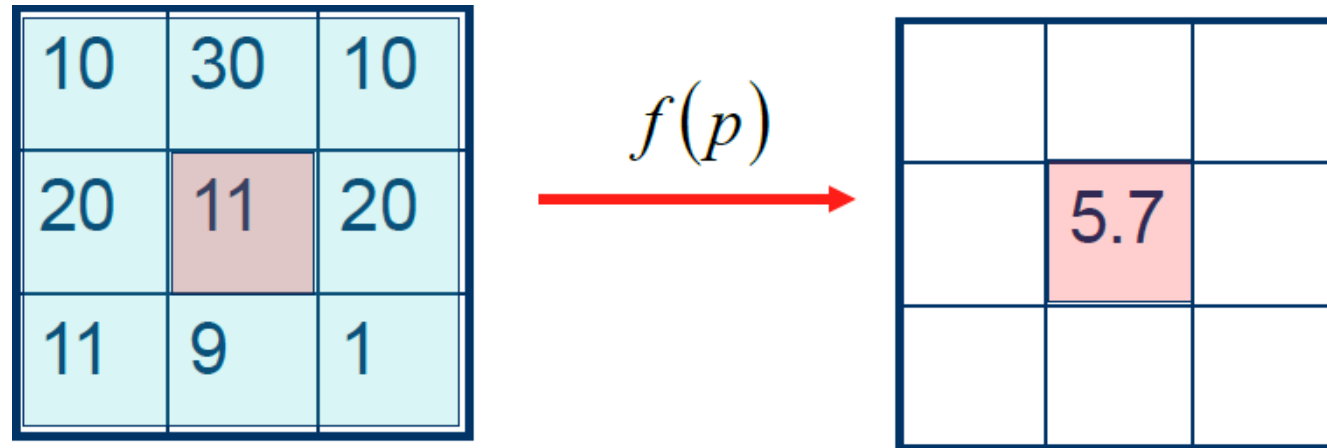
[] [] []

Image filtering

- Image filtering: compute function of local neighborhood at each position
- Enhance images
 - Denoise, resize, increase contrast, etc.
- Extract information from images
 - Texture, edges, distinctive points, etc.
- Detect patterns
 - Template matching

Filtering

- Modify pixels based on some function of neighborhood



Filtering

- Output is linear combination of the neighborhood pixels

1	3	0
2	10	2
4	1	1

 \otimes

1	0	-1
1	0.1	-1
1	0	-1

 $=$

	5	

Image

Kernel

Filter Output

Derivatives and Average

- **Derivative:** rate of change
 - Speed is a rate of change of a distance, $X=V.t$
 - Acceleration is a rate of change of speed, $V=a.t$
- **Average:** mean
 - Dividing the sum of N values by N

Derivative

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x$$

$$y = x^2 + x^4$$

$$\frac{dy}{dx} = 2x + 4x^3$$

$$y = \sin x + e^{-x}$$

$$\frac{dy}{dx} = \cos x + (-1)e^{-x}$$

Discrete Derivative

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x)$$

$$\frac{df}{dx} = \frac{f(x) - f(x - 1)}{1} = f'(x)$$

$$\frac{df}{dx} = f(x) - f(x - 1) = f'(x)$$

Discrete Derivative / Finite Difference

$$\frac{df}{dx} = \frac{f(x) - f(x-1)}{1} = f'(x) \quad \text{Backward difference}$$

$$\frac{df}{dx} = \frac{f(x) - f(x+1)}{-1} = f'(x) \quad \text{Forward difference}$$

$$\frac{df}{dx} = \frac{f(x+1) - f(x-1)}{2} = f'(x) \quad \text{Central difference}$$

Example: Finite Difference

$$f(x) = 10 \quad 15 \quad 10 \quad 10 \quad 25 \quad 20 \quad 20 \quad 20$$

$$f'(x) = 0 \quad 5 \quad -5 \quad 0 \quad 15 \quad -5 \quad 0 \quad 0$$

$$f''(x) = 0 \quad 5 \quad 10 \quad 5 \quad 15 \quad -20 \quad 5 \quad 0$$

Derivative Masks

Backward difference $[-1 \quad 1]$

Forward difference $[1 \quad -1]$

Central difference $[-1 \quad 0 \quad 1]$

Derivative in 2-D

Given function

$$f(x, y)$$

Gradient vector

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} \\ \frac{\partial f(x, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

Gradient magnitude

$$|\nabla f(x, y)| = \sqrt{f_x^2 + f_y^2}$$

Gradient direction

$$\theta = \tan^{-1} \frac{f_x}{f_y}$$

Derivative of Images

Derivative masks

$$f_x \Rightarrow \frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad f_y \Rightarrow \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix}$$

The image matrix I is a 5x5 grid of values. A red box highlights the first three columns (values 10, 10, 20), and a blue box highlights the last three columns (values 10, 20, 20). The overlap of these boxes covers the central 3x3 region of the matrix.

$$I_x = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

The derivative image I_x is a 5x5 grid of values. A red box highlights the value 10 at row 2, column 2, and a blue box highlights the value 10 at row 2, column 3. These boxes are positioned over the central 3x3 region of the matrix.

Derivative of Images

Derivative masks

$$f_x \Rightarrow \frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad f_y \Rightarrow \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} \quad I_y = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Averages

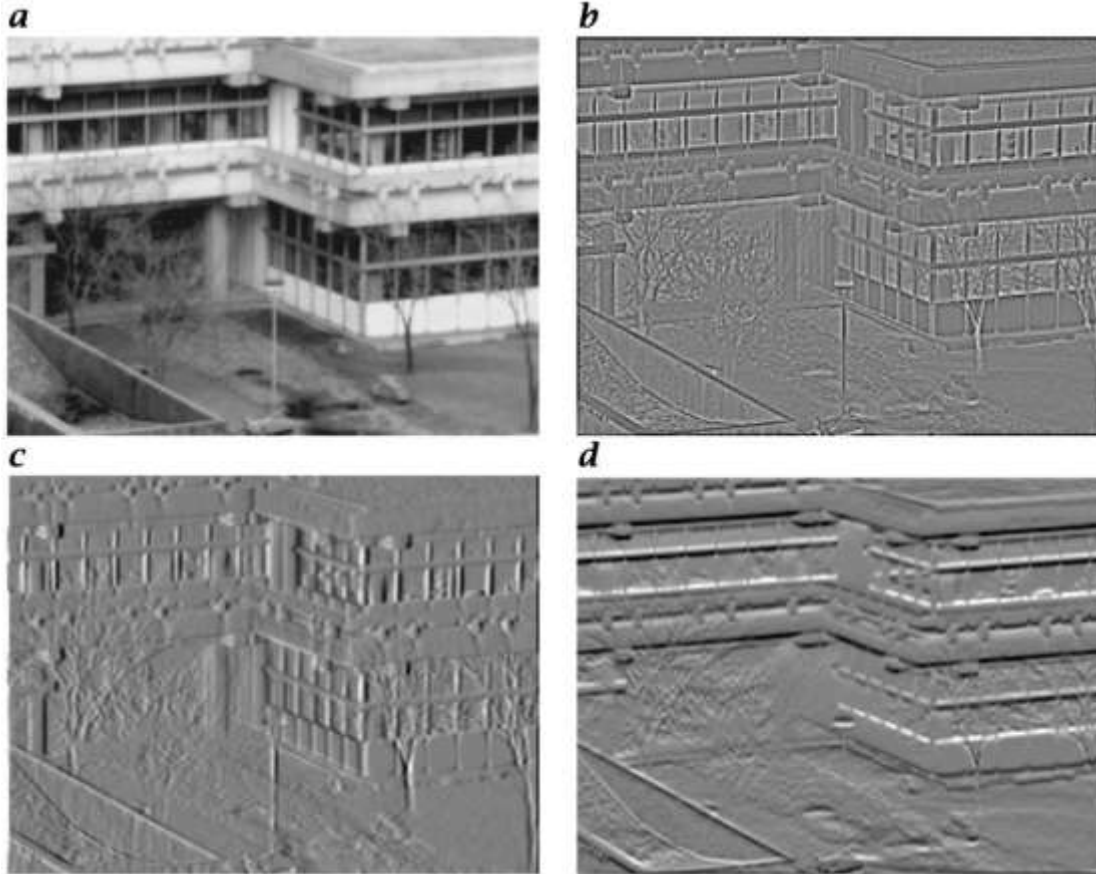
- Mean

$$I = \frac{I_1 + I_2 + \dots + I_n}{n} = \frac{\sum_{i=1}^n I_i}{n}$$

- Weighted mean

$$I = \frac{w_1 I_1 + w_2 I_2 + \dots + w_n I_n}{n} = \frac{\sum_{i=1}^n w_i I_i}{n}$$

Example



- a. Original image
- b. Laplacian operator
- c. Horizontal derivative
- d. Vertical derivative

Correlation (linear relationship)

$$f \otimes h = \sum \sum f(k$$

f = Image

h = Kernel

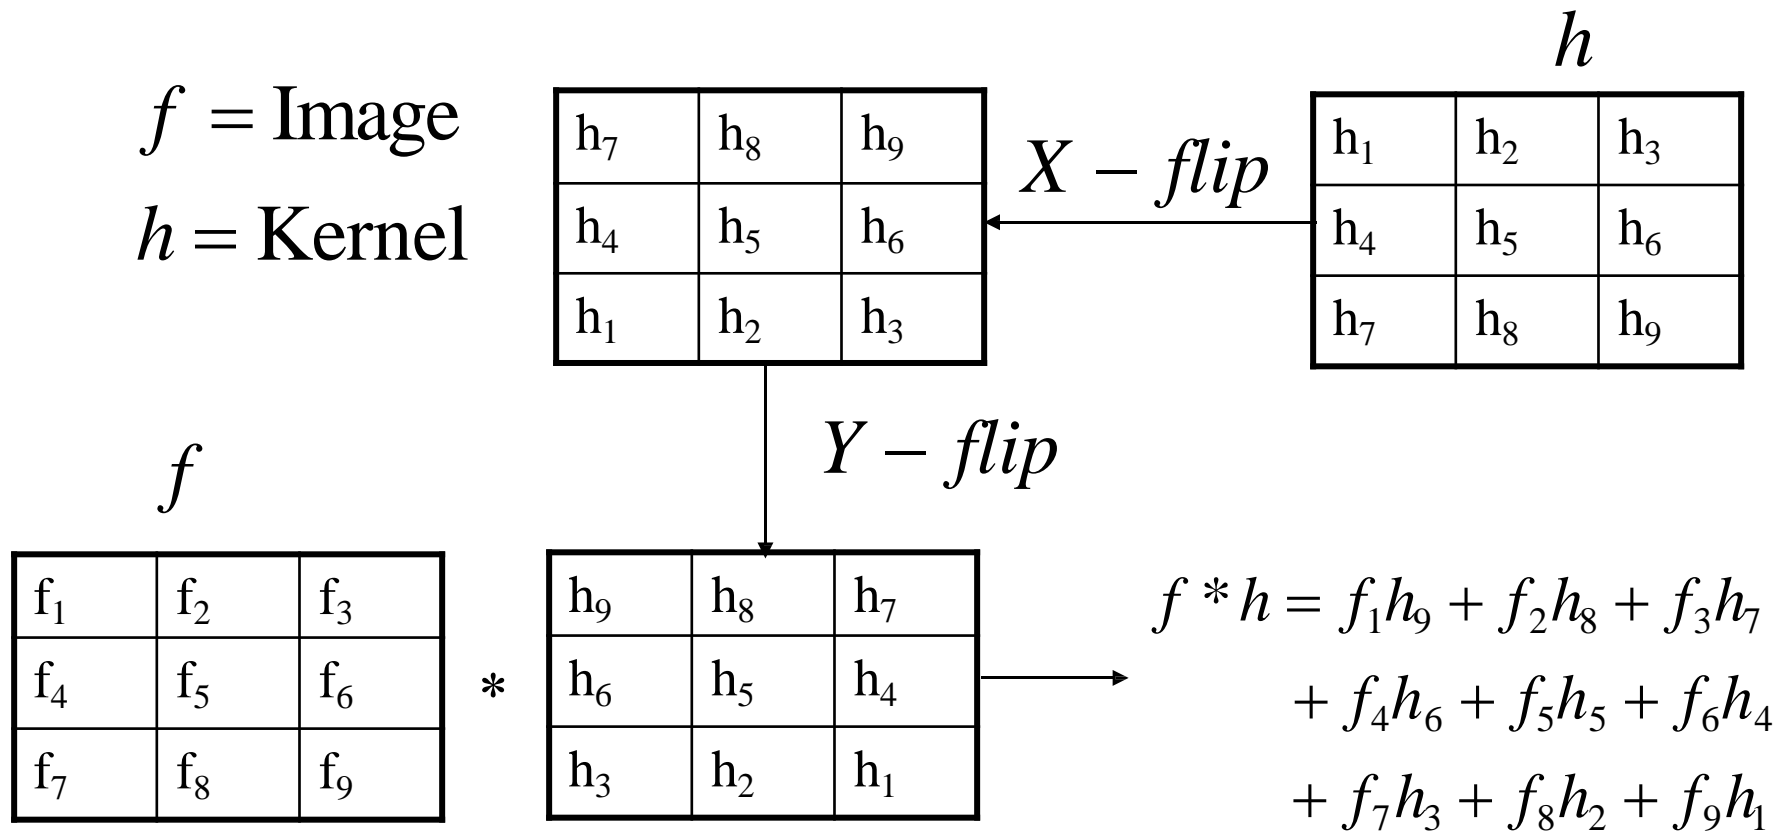
$$\begin{array}{c} f \\ \begin{array}{|c|c|c|} \hline f_1 & f_2 & f_3 \\ \hline f_4 & f_5 & f_6 \\ \hline f_7 & f_8 & f_9 \\ \hline \end{array} \end{array} \otimes \begin{array}{c} h \\ \begin{array}{|c|c|c|} \hline h_1 & h_2 & h_3 \\ \hline h_4 & h_5 & h_6 \\ \hline h_7 & h_8 & h_9 \\ \hline \end{array} \end{array} \rightarrow \begin{aligned} f \otimes h &= f_1h_1 + f_2h_2 + f_3h_3 \\ &+ f_4h_4 + f_5h_5 + f_6h_6 \\ &+ f_7h_7 + f_8h_8 + f_9h_9 \end{aligned}$$

Convolution

$$f * h = \sum \sum f(k, l)$$

f = Image

h = Kernel



Convolution

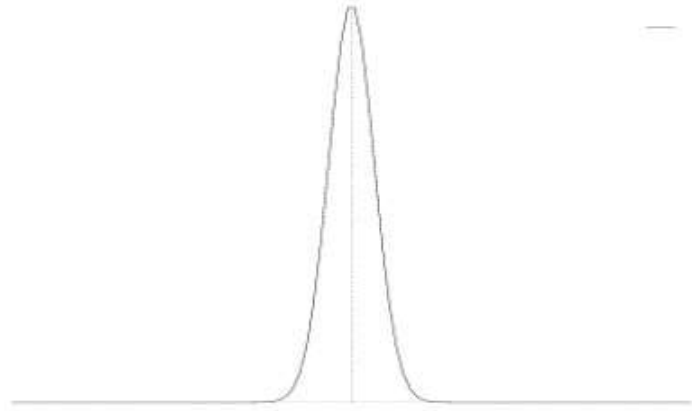
- Convolution is **associative**

$$F * (G * I) = ($$

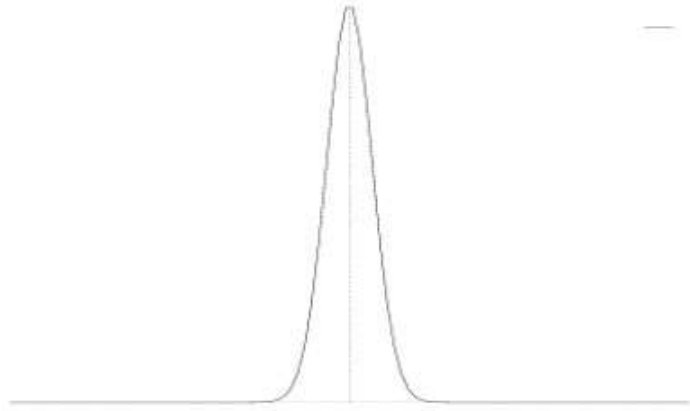
Correlation and Convolution

- **Convolution** is a filtering operation
 - expresses **the amount of overlap** of one function as it is shifted over another function
- **Correlation** compares the similarity of two sets of data
 - relatedness of the signals!

Gaussian filter

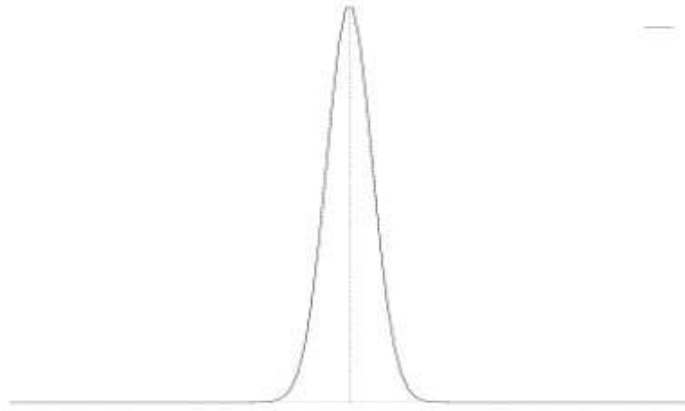


Gaussian filter

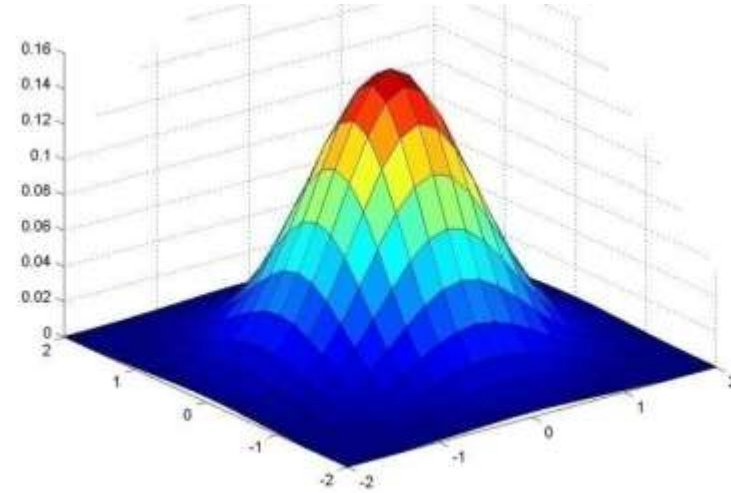


$$g(x) = e^{\frac{-x^2}{2\sigma^2}}$$

Gaussian filter

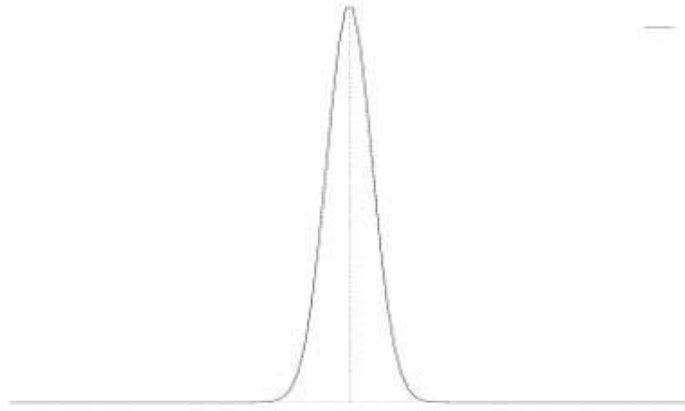


$$g(x) = e^{\frac{-x^2}{2\sigma^2}}$$

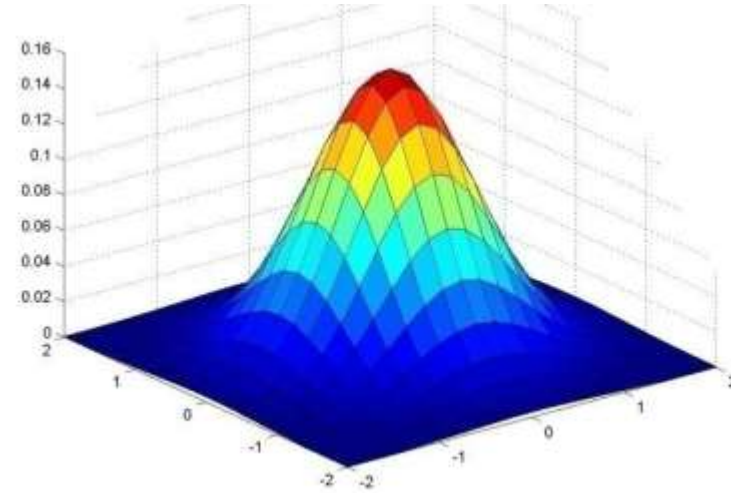


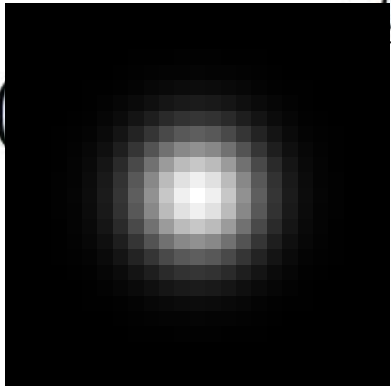
$$g(x, y) = e^{\frac{-(x^2 + y^2)}{2\sigma^2}}$$

Gaussian filter



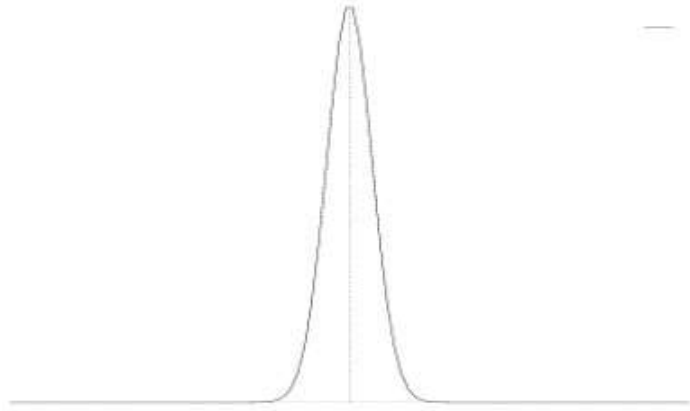
$$g(x) = e^{\frac{-x^2}{2\sigma^2}}$$



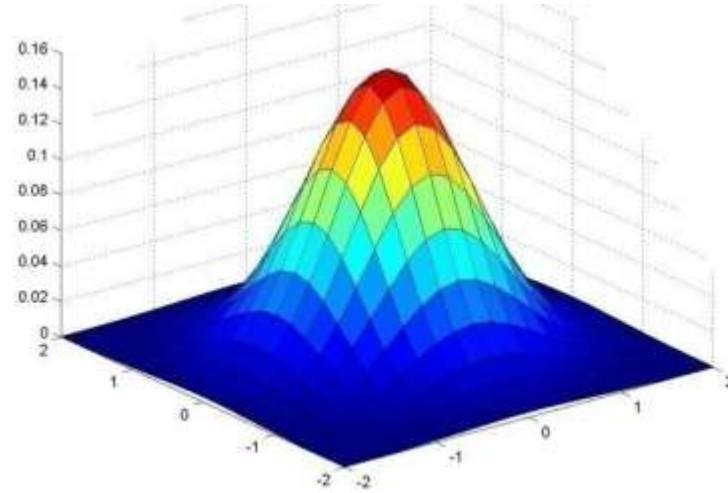
$$g(x, y) = e^{\frac{-(x^2 + y^2)}{2\sigma^2}}$$


A 2D grayscale plot of a Gaussian function, showing a circular, blurred spot centered at the origin. The plot is square-shaped with a black background and a white center.

Gaussian filter

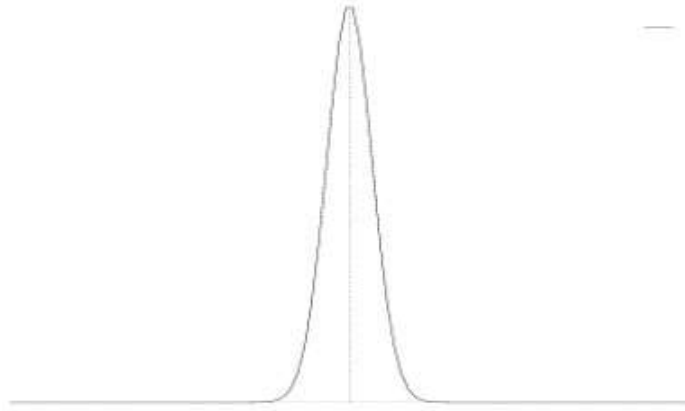


$$g(x) = e^{\frac{-x^2}{2\sigma^2}}$$

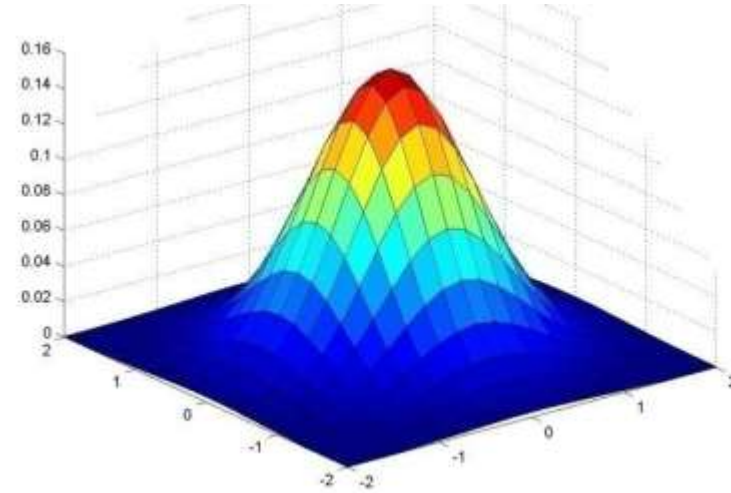


$$g(x, y) = e^{\frac{-(x^2 + y^2)}{2\sigma^2}}$$

Gaussian filter



$$g(x) = e^{\frac{-x^2}{2\sigma^2}}$$

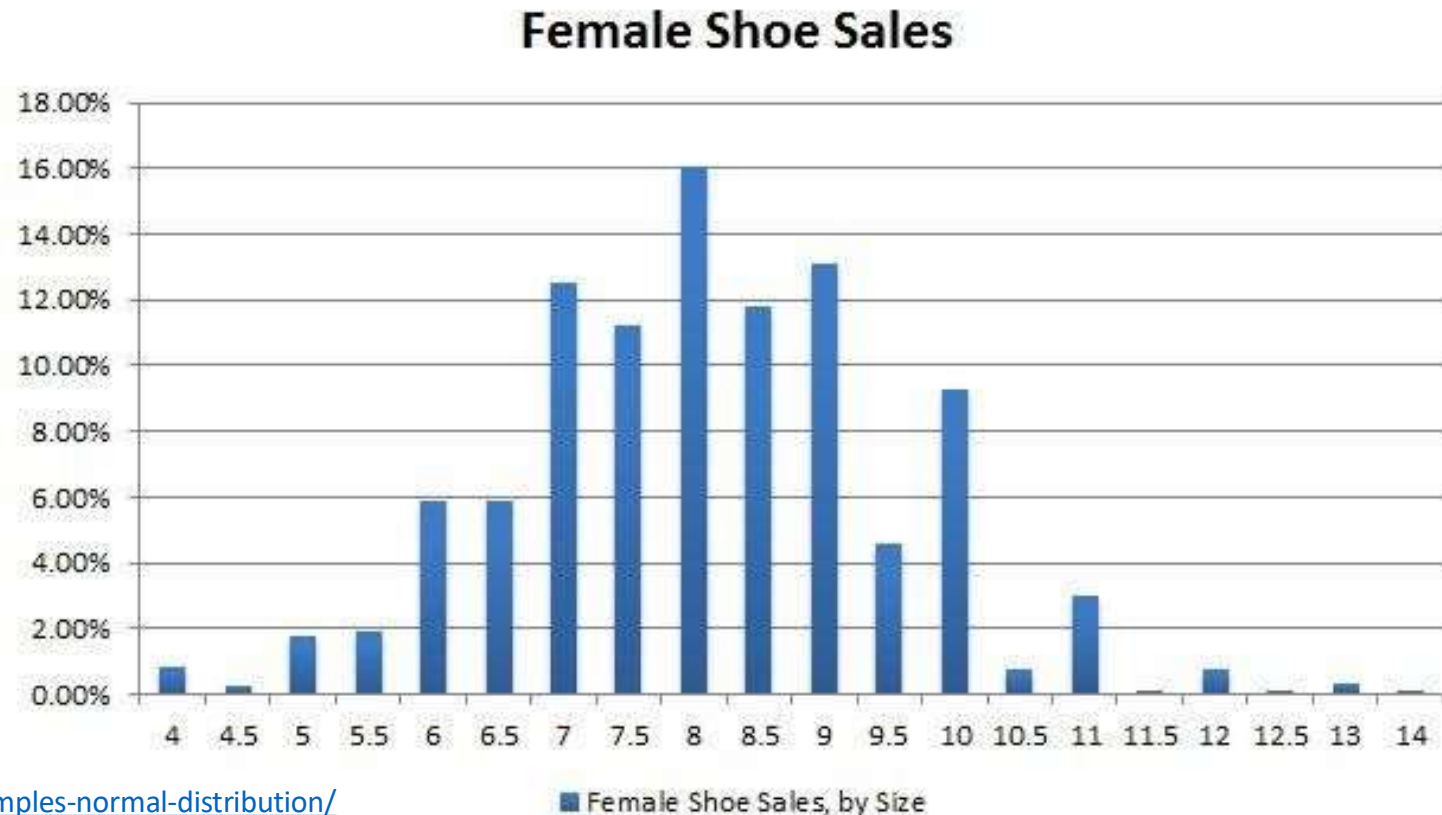


$$g(x, y) = e^{\frac{-(x^2 + y^2)}{2\sigma^2}}$$

$$g(x) = \begin{bmatrix} .011 & .13 & .6 & 1 & .6 & .13 & .011 \end{bmatrix}$$

Gaussian filter - properties

- Most common natural model

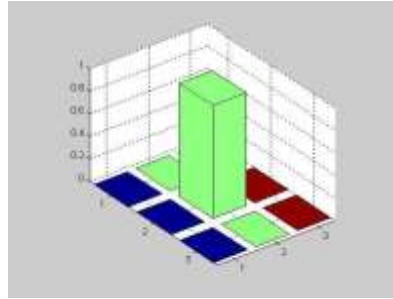


<https://studiousguy.com/real-life-examples-normal-distribution/>

Gaussian filter - properties

- Most common natural model
- Smooth function, it has infinite number of derivatives
- It is Symmetric
- Fourier Transform of Gaussian is Gaussian.
- Convolution of a Gaussian with itself is a Gaussian.
- Gaussian is separable; 2D convolution can be performed by two 1-D convolutions
- There are cells in eye that perform Gaussian filtering.

Filtering Examples - 1



*

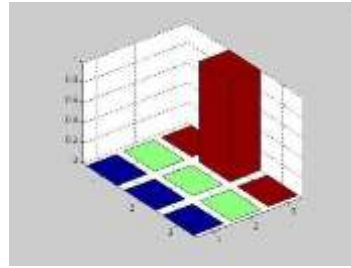
0	0	0
0	1	0
0	0	0



Filtering Examples - 2



*

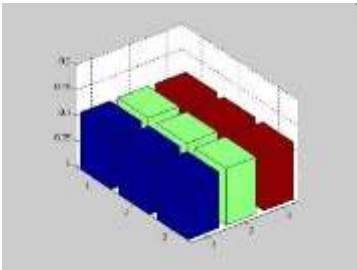


0	0	0
1	0	0
0	0	0


=



Filtering Examples - 3



A 3D bar chart representing a 3x3x3 kernel. The front face (z=0) is blue, the middle face (z=1) is green, and the back face (z=2) is red. The axes are labeled 1, 2, 3.




A grayscale image of a woman wearing a hat, which is the input to the filtering operation.

$$\ast \frac{1}{9}$$

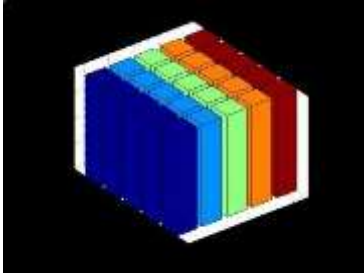
1	1	1
1	1	1
1	1	1

A 3x3 grid of ones, representing the averaging kernel.

$$=$$


A grayscale image of the same woman wearing a hat, which is the output of the filtering operation. The image is slightly blurred compared to the input.

Filtering Examples - 4




A 3D visualization of a 5x5x5 volume, likely representing a 3D filter kernel, with colored layers (blue, green, yellow, orange, red) visible.

A grayscale image of a woman wearing a hat, representing the input image.

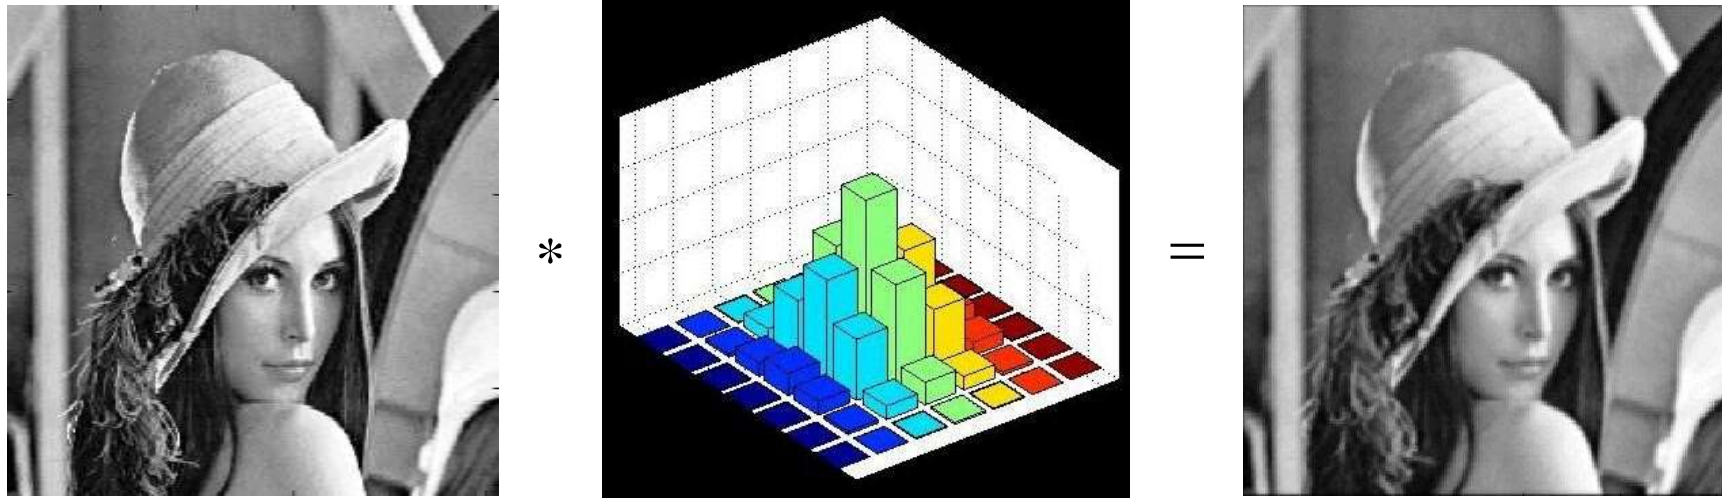
$$* \frac{1}{25}$$

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

=



Filtering Examples - 5



Gaussian Smoothing

Filtering Examples - 6



Gaussian Smoothing



Smoothing by Averaging

Filtering Examples - 7



After additive
Gaussian Noise



After Averaging



After Gaussian Smoothing

Questions?

Sources for this lecture include materials from works by Mubarak Shah, S. Seitz, James Tompkin and Ulas Bagci